TRAIL OF BITS

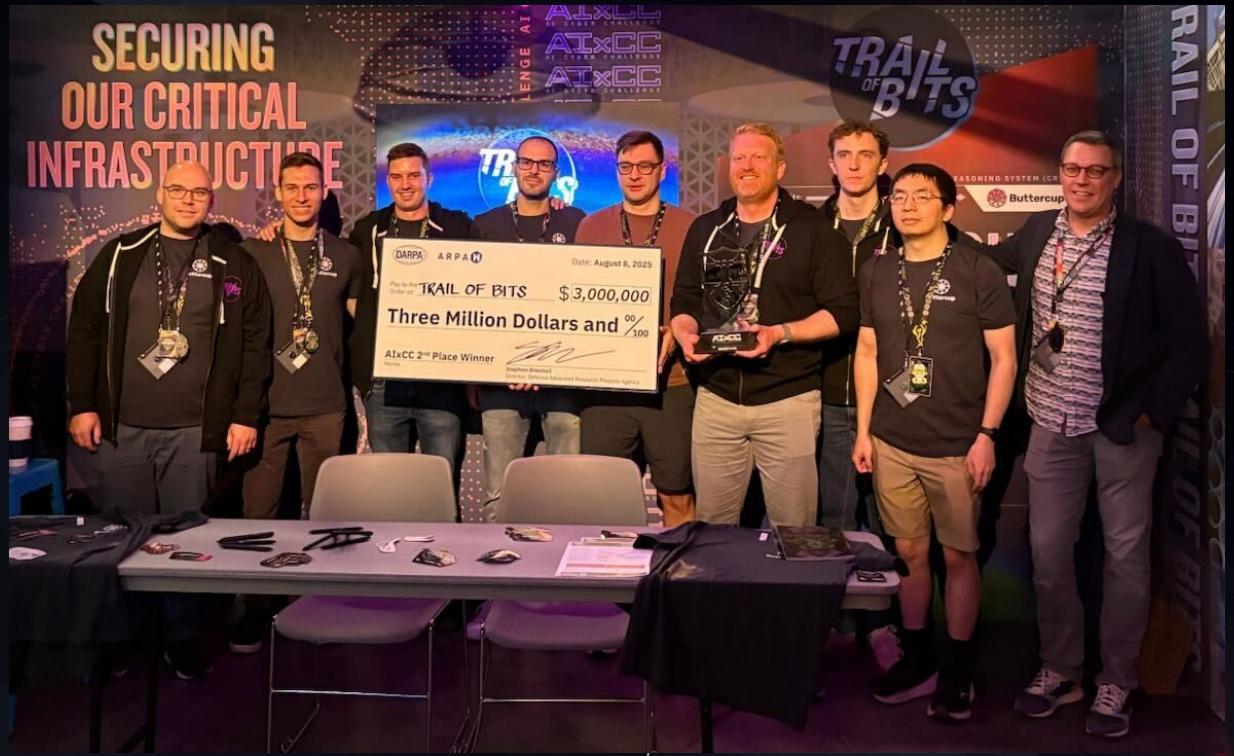

Buttercup: The Future of Trail of Bits' Solution to the DARPA's AI Cyber Challenge

2024 CrowdStrike IT outage

What does a bug look like?

```
public class MessageProcessor {
    public byte[] processMessage(int messageLength, int headerSize) {
        // Calculate total size needed
        int totalSize = messageLength + headerSize + 1024; // extra padding

        // Allocate buffer
        byte[] buffer = new byte[totalSize];

        // Process message...
        return buffer;
    }
}
```

What does a patch look like?

```java
public class MessageProcessor {
    private static final int MAX_MESSAGE_SIZE = 10_000_000; // 10MB limit

    public byte[] processMessage(int messageLength, int headerSize) {
        // Validate inputs first
        if (messageLength < 0 || messageLength > MAX_MESSAGE_SIZE) {
            throw new IllegalArgumentException(
                "Invalid message length: " + messageLength);
        }
        if (headerSize < 0 || headerSize > 1024) {
            throw new IllegalArgumentException(
                "Invalid header size: " + headerSize);
        }

        // Check for overflow before doing arithmetic
        if (messageLength > MAX_MESSAGE_SIZE - headerSize - 1024) {
            throw new IllegalArgumentException(
                "Message too large: would cause overflow" );
        }

        int totalSize = messageLength + headerSize + 1024;
        byte[] buffer = new byte[totalSize];

        return buffer;
    }
}
```

Key Scoring Components

Points Awarded For:

- **Patches** - Worth the most points
- **Proof of Vulnerabilities** - Worth moderate points
- **Bundled submissions** - Extra points when POV + patch submitted
- **Static analysis reports** - Minor points (mainly useful when bundled)

Scoring Modifiers:

- **Speed bonus** - Faster submissions earn more points
- **Accuracy multiplier** - Incorrect/duplicate submissions reduce your score multiplier
- **Baseline comparison** - Must outperform state-of-the-art baseline systems

Penalties:

- Duplicate POVs or patches reduce accuracy multiplier
- Incorrect patches (that don't fix the bug or break functionality) harm accuracy multiplier
- Submitting patches that get "clobbered" by later POVs loses points

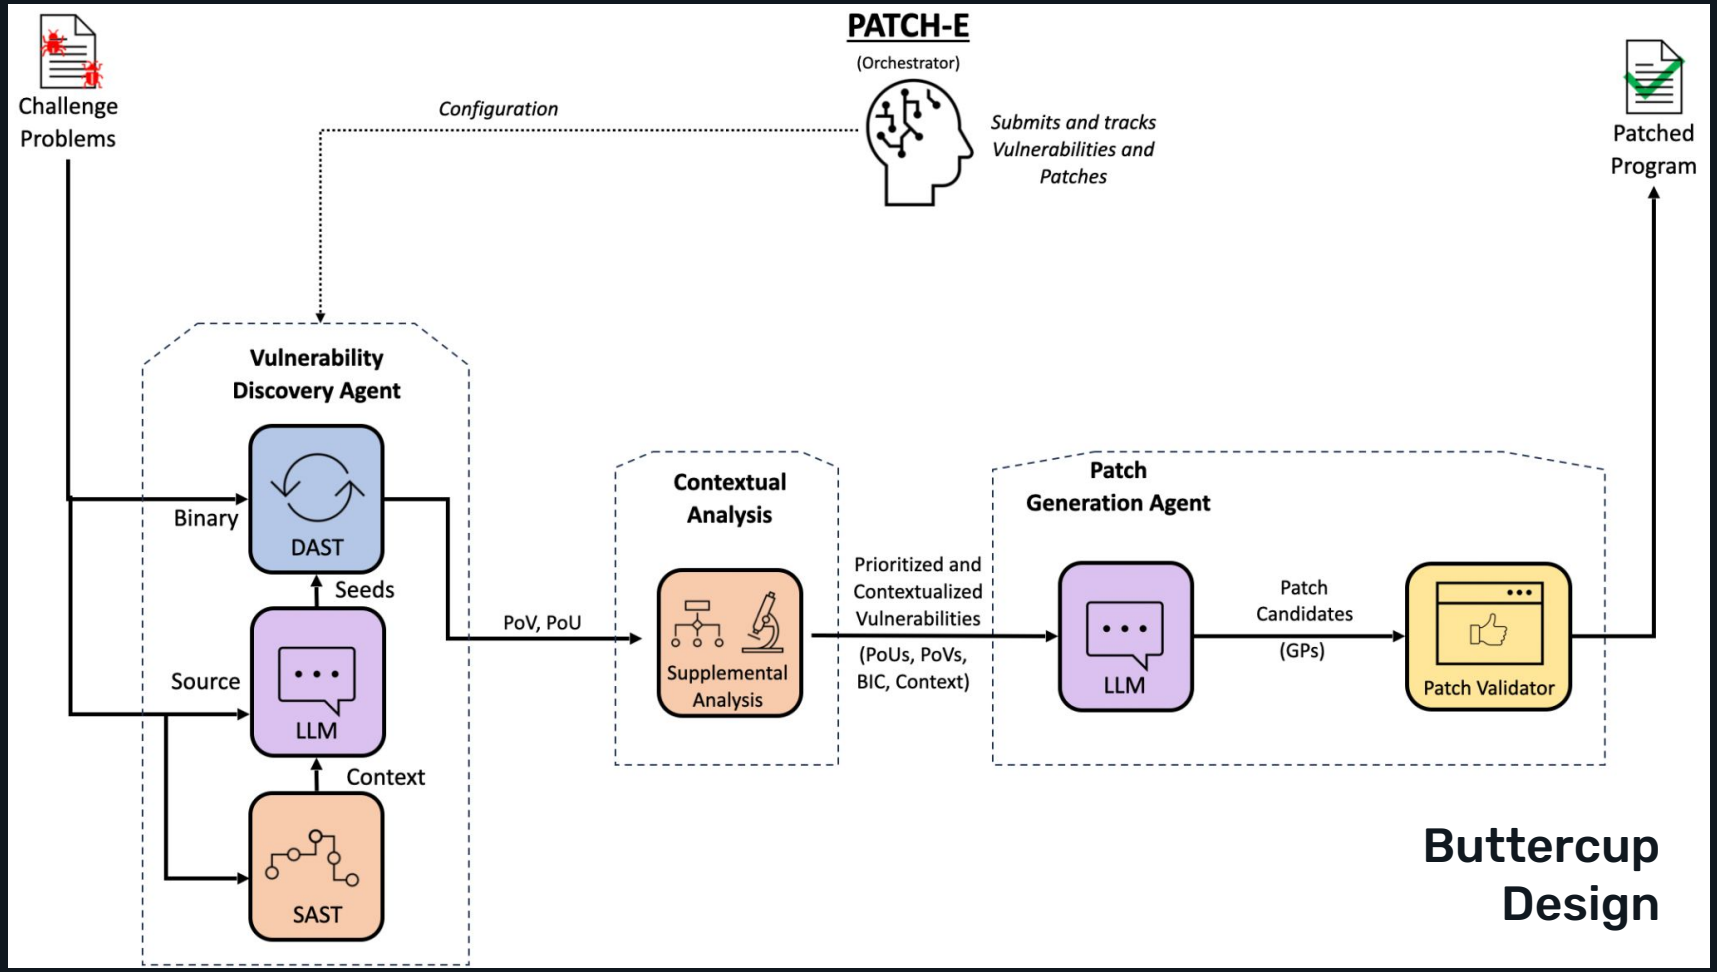
PATCH-E
(Orchestrator)
Challenge Problems
Configuration
Submits and tracks Vulnerabilities and Patches
Patched Program
Vulnerability Discovery Agent
Binary
DAST
Seeds
Source
LLM
Context
SAST
PoV, PoU
Contextual Analysis
Supplemental Analysis
Prioritized and Contextualized Vulnerabilities
(PoUs, PoVs, BIC, Context)
Patch Generation Agent
LLM
Patch Candidates
(GPs)
Patch Validator
Buttercup Design

Team	LLM spend	Compute spend	Total spend	Cost per point
Team Atlanta	$29.4k	$73.9k	$103.3k	$263
Trail of Bits	**$21.1k**	**$18.5k**	**$39.6k**	**$181**
Theori	$11.5k	$20.3k	$31.8k	$151
fuzzing_brain	$12.2k	$63.2k	$75.4k	$490
Shellphish	$2.9k	$54.9k	$57.8k	$425
42-b3yond-6ug	$1.1k	$38.7k	$39.8k	$379
LACROSSE	$631	$7.1k	$7.2k	$751

Results!

2nd in LLM Spend
6th in Compute Spend
2nd in $ per Point

The Future of AI in Cybersecurity

- Determining exploitable vulns
- Binary reverse engineering
- Formal methods reasoning
- No more malware C&Cs
- 🤫 No more SBIR Phase I's

trent@trailofbits.com
www.trailofbits.com/buttercup

Our work, blog.trailofbits.com
Our code, github.com/trailofbits
Our socials, @trailofbits on X