



DECEMBER 10-11, 2025

EXCEL LONDON / UNITED KINGDOM

Cut To The QUIC

Slashing QUIC's Performance With A Hash DoS

Paul Bottinelli

Principal Security Engineer, Cryptography @ Trail of Bits



Paul Bottinelli

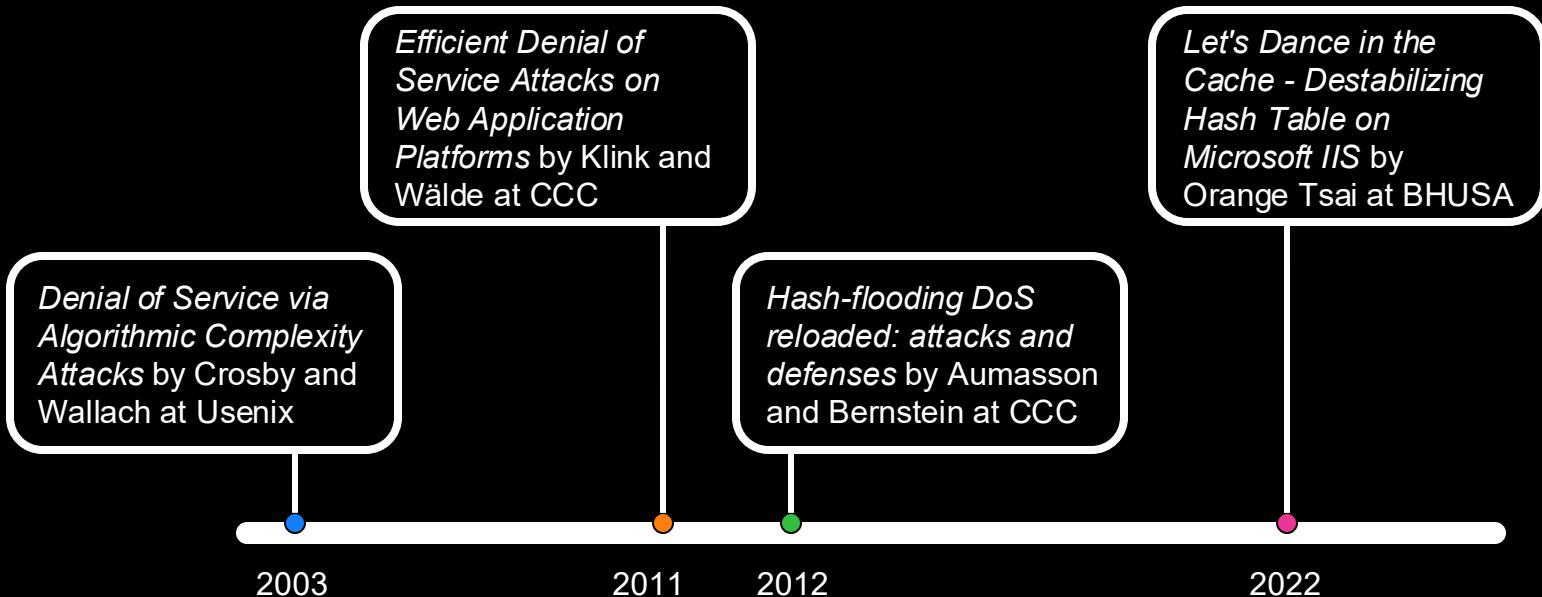
- Cryptography consultant



15 years ago...



Hash DoS – A Long History



Hash DoS Ingredients



Hash Tables

Hash Tables

```
{  
    "Alice": "10.0.0.1",  
    "Bob": "10.0.0.2",  
    "Claire": "10.0.0.3",  
    "Dan": "10.0.0.4"  
}
```

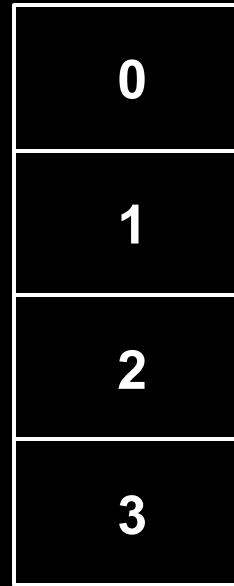
Insertion and lookup
complexity

$O(1)^*$

Insertion

```
function insert(key, value)
    index = hash(key) % len(array)
    array[index] = (key, value)

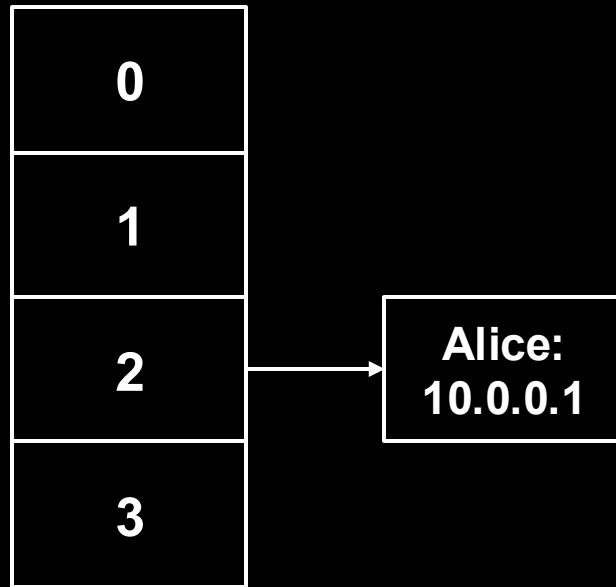
> insert("Alice", "10.0.0.1")
    hash("Alice") = 6
    index = 6 % 4 = 2
```



Insertion

```
function insert(key, value)
    index = hash(key) % len(array)
    array[index] = (key, value)

> insert("Alice", "10.0.0.1")
    hash("Alice") = 6
    index = 6 % 4 = 2
    array[2] = ("Alice", "10.0.0.1")
```

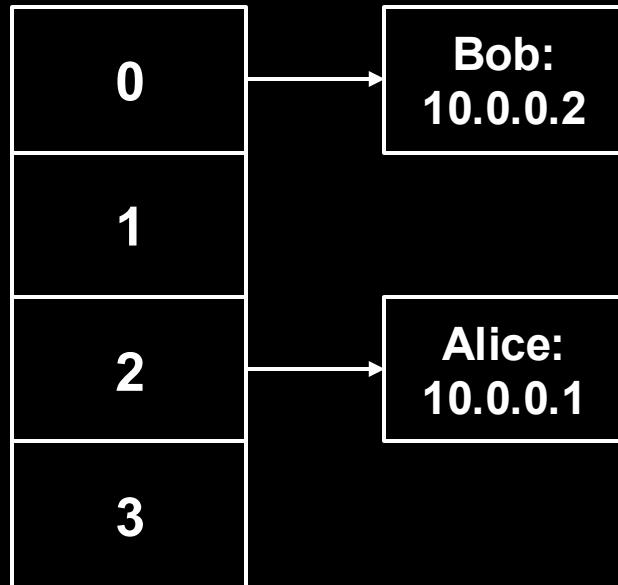


Insertion

```
function insert(key, value)
    index = hash(key) % len(array)
    array[index] = (key, value)

> insert("Alice", "10.0.0.1")
[SUCCESS] inserted at index 2

> insert("Bob", "10.0.0.2")
[SUCCESS] inserted at index 0
```



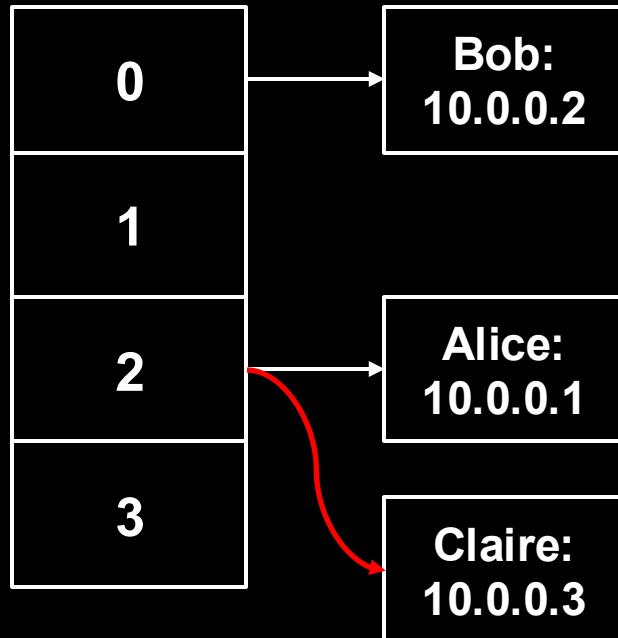
Insertion

```
function insert(key, value)
    index = hash(key) % len(array)
    array[index] = (key, value)

> insert("Alice", "10.0.0.1")
[SUCCESS] inserted at index 2

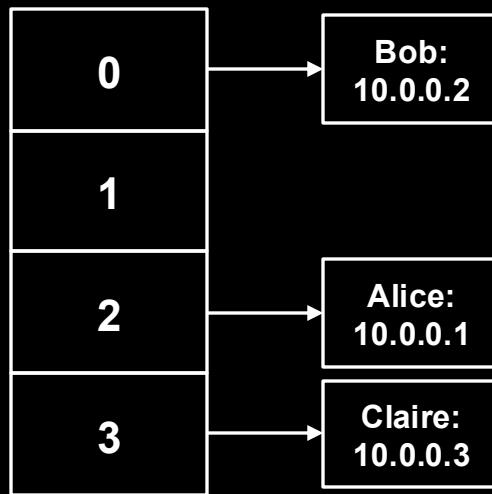
> insert("Bob", "10.0.0.2")
[SUCCESS] inserted at index 0

> insert("Claire", "10.0.0.3")
[ERROR] index 2 full!
```

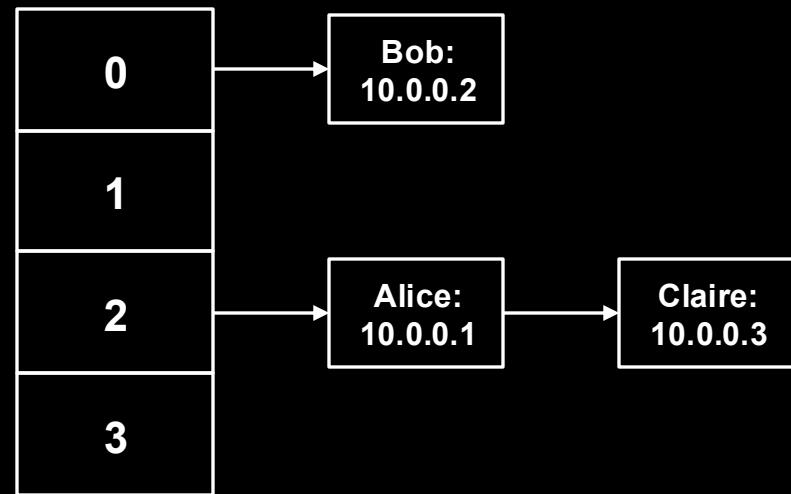


Collision Resolution Strategies

Open addressing



Separate chaining



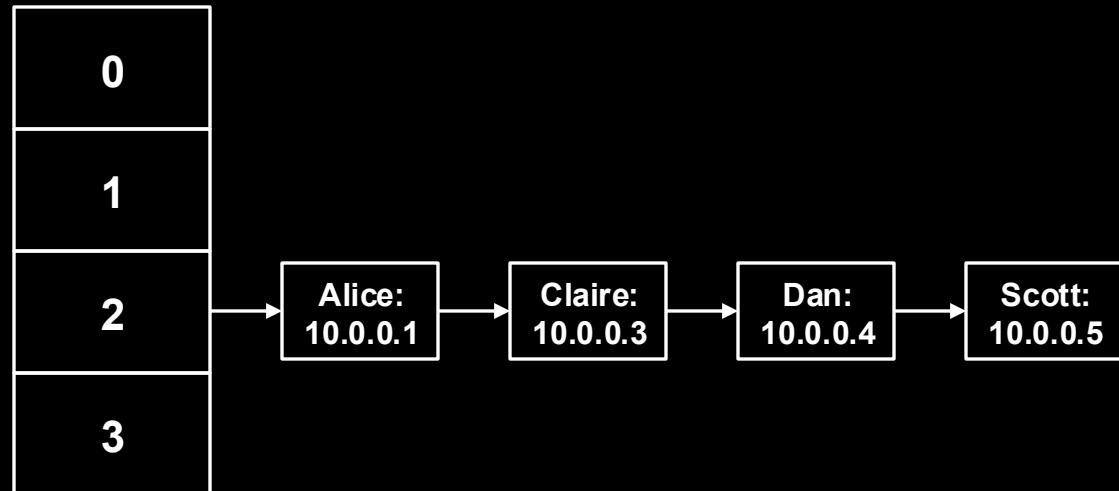
Hash Denial of Service

Lookup complexity of colliding elements

$$O(n)$$

Amortized complexity with n elements

$$O(n^2)$$



Hash DoS are **algorithmic complexity** attacks.

Hash DoS Ingredients



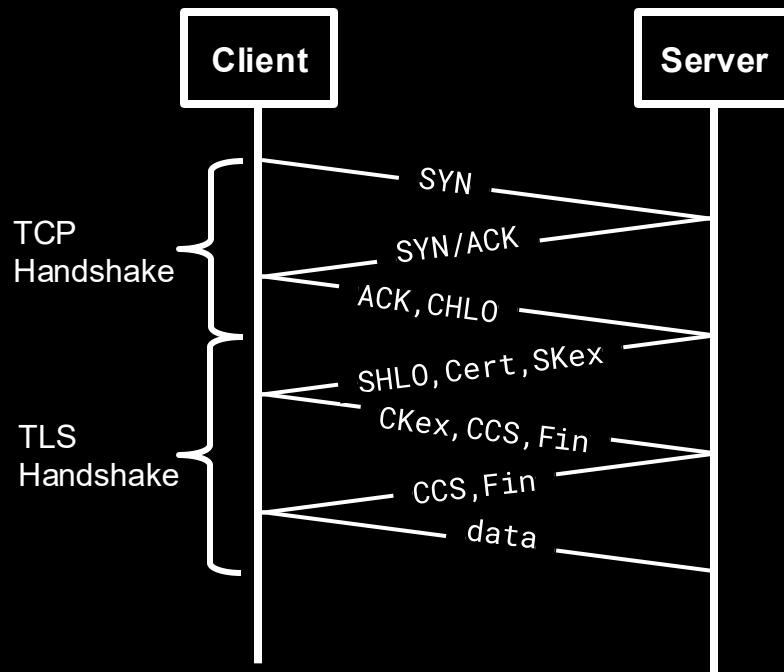
QUIC (quickly)

QUIC

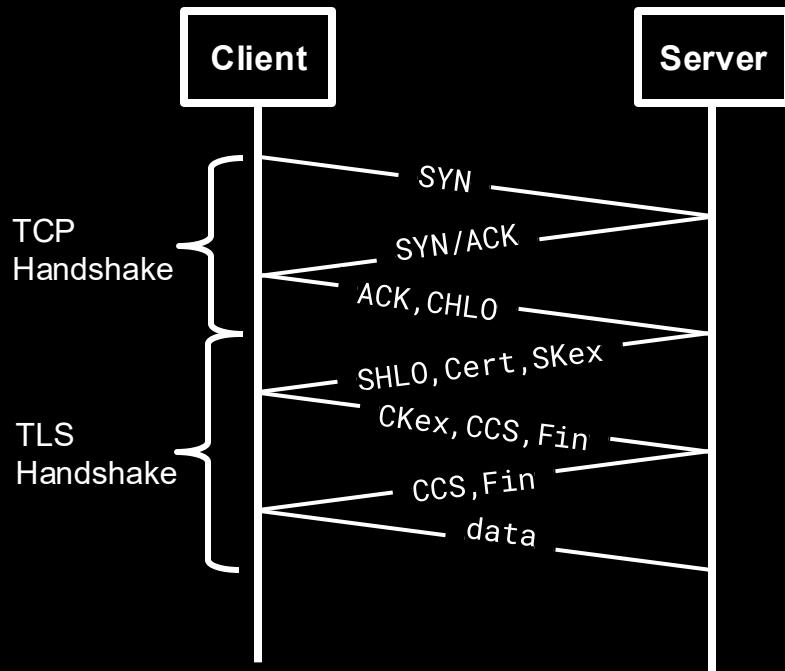
- Transport protocol originally designed by Google in 2012
- Formalized under [RFC 9000](#) (and [8999](#), [9001](#), and [9002](#))
- Backbone of HTTP/3
- Improve performance of web applications and reduce network latency with UDP
- Leverages TLS 1.3 handshake



TCP+TLS

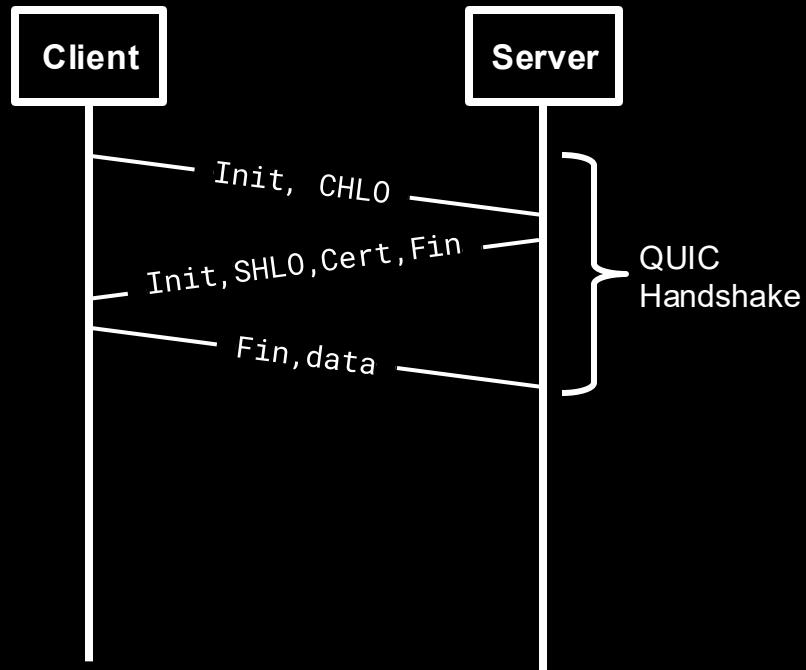


TCP+TLS



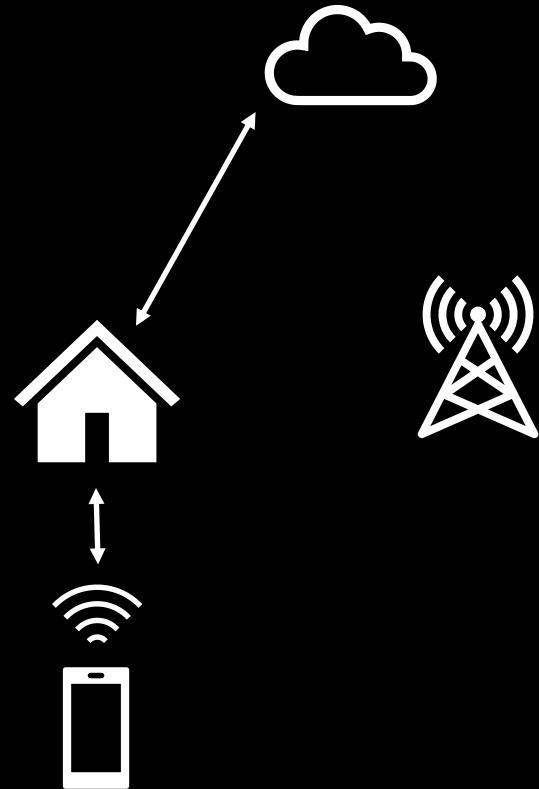
VS

QUIC



QUIC Connections

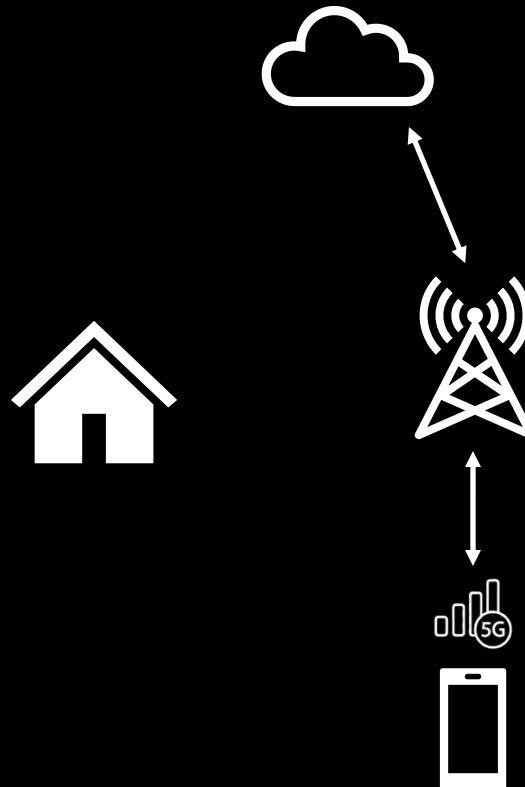
Improve performance during
network-switching events



QUIC Connections

Improve performance during
network-switching events

- Backend server keeps track of connections using **Connection IDs**
- Regardless of any changes in the source IP address
- Upon switching networks connection is re-established by sending a packet containing CID

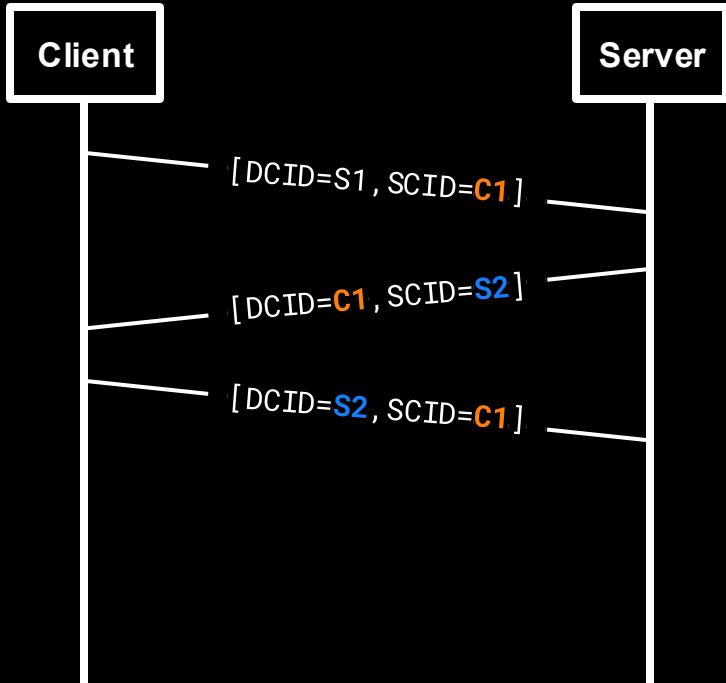


QUIC Connections

- A connection is *uniquely* identified by two connection IDs, the Source CID (SCID) and the Destination CID (DCID)



- The client selects Source Connection ID using an implementation-specific (...) method ([RFC 9000, Section 5.1](#))
- Natural for the server to store them in a hash table indexed by CID
- SCID is attacker-controlled → Hash DoS !



Hash DoS Ingredients



Hash DoS in QUIC

Modern programming languages protect against Hash DoS, see Rust's [HashMap](#):

By default, `HashMap` uses a hashing algorithm selected to provide resistance against HashDoS attacks. The algorithm is randomly seeded,

How easy is it to find vulnerable implementations?

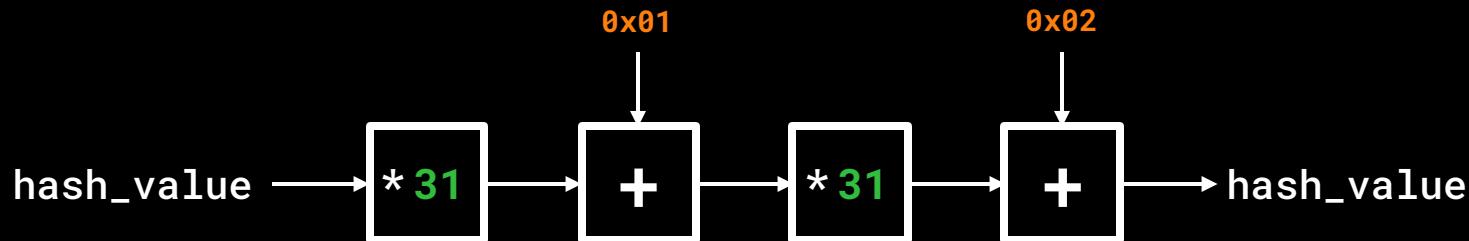
Multiplicative Hash in xquic

```
uint64_t xqc_hash_string(const u_char *data, size_t len){  
    uint64_t hash_value = 0;  
    for (size_t i = 0; i < len; ++i) {  
        hash_value = hash_value * 31 + data[i];  
    }  
    return hash_value;  
}
```

Multiplicative Hash in xquic

```
uint64_t xqc_hash_string(const u_char *data  size_t len){  
    uint64_t hash_value = 0;  
    for (size_t i = 0; i < len; ++i) {  
        hash_value = hash_value * 31 + data[i];  
    }  
    return hash_value;  
}
```

> xqc_hash_string({0x01, 0x02}, 2)



Multiplicative Hash in xquic

```
uint64_t xqc_hash_string(const u_char *data  size_t len){  
    uint64_t hash_value = 0;  
    for (size_t i = 0; i < len; ++i) {  
        hash_value = hash_value * 31 + data[i];  
    }  
    return hash_value;  
}
```

> xqc_hash_string({0x01, 0x02}, 2)

hash_value = 33

Computing Collisions in xquic

$$\text{hash}(0x\text{XXYY}) = 31 * \text{XX} + \text{YY}$$

$$\begin{aligned}\text{hash}(0x00ff) &= 31 * 0 + 255 = 255 \\ &\quad 31 * 1 + 224 = 255 = \text{hash}(0x01e0) \\ &\quad 31 * 2 + 193 = 255 = \text{hash}(0x02c1) \\ &\quad \dots \\ &\quad 31 * 8 + 7 = 255 = \text{hash}(0x0807)\end{aligned}\right.$$

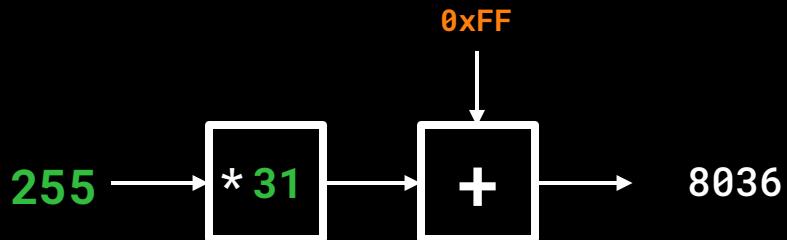
The nine 2-byte arrays all hash to the same value

→ Equivalent Substring Attack!

Equivalent Substring Attack in xquic

```
uint64_t xqc_hash_string(const u_char *data  size_t len){  
    uint64_t hash_value = 0;  
    for (size_t i = 0; i < len; ++i) {  
        hash_value = hash_value * 31 + data[i];  
    }  
    return hash_value;  
}
```

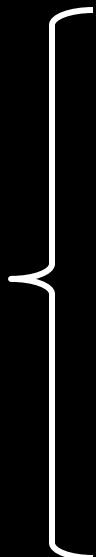
```
> xqc_hash_string({0x00,0xFF}, 2)  
> xqc_hash_string({0x00,0xFF,0xFF}, 3)  
= 8036  
> xqc_hash_string({0x01,0xE0,0xFF}, 3)  
= 8036
```



Appending the same suffix to colliding arrays maintains the hash collision!

Equivalent Substring Attack on xquic

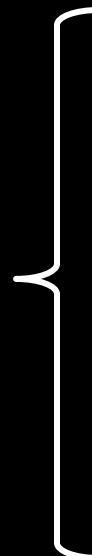
The nine 2-byte arrays all hash to the same value



hash(0x00ff)
hash(0x01e0)
hash(0x02c1)
hash(0x03a2)
hash(0x0483)
hash(0x0564)
hash(0x0645)
hash(0x0726)
hash(0x0807)

Equivalent Substring Attack on xquic

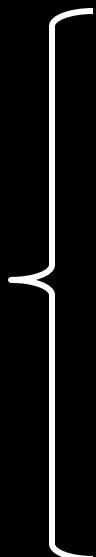
The nine 2-byte arrays all hash to the same value



```
hash(0x00ff      )  
hash(0x01e0      )  
hash(0x02c1      )  
hash(0x03a2      )  
hash(0x0483      )  
hash(0x0564      )  
hash(0x0645      )  
hash(0x0726      )  
hash(0x0807      )
```

Equivalent Substring Attack on xquic

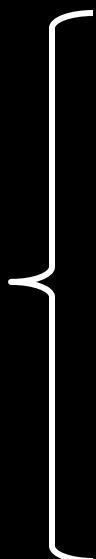
The nine 4-byte arrays all hash to the same value



```
hash(0x00ff abcd )  
hash(0x01e0 abcd )  
hash(0x02c1 abcd )  
hash(0x03a2 abcd )  
hash(0x0483 abcd )  
hash(0x0564 abcd )  
hash(0x0645 abcd )  
hash(0x0726 abcd )  
hash(0x0807 abcd )
```

Equivalent Substring Attack on xquic

The nine 4-byte arrays all hash to the same value



```
hash(0x00ff 00ff )  
hash(0x01e0 00ff )  
hash(0x02c1 00ff )  
hash(0x03a2 00ff )  
hash(0x0483 00ff )  
hash(0x0564 00ff )  
hash(0x0645 00ff )  
hash(0x0726 00ff )  
hash(0x0807 00ff )
```

Equivalent Substring Attack on xquic

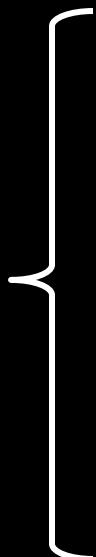
The nine 4-byte arrays all hash to the same value



```
hash(0x00ff 00ff )  
hash(0x01e0 01e0 )  
hash(0x02c1 02c1 )  
hash(0x03a2 03a2 )  
hash(0x0483 0483 )  
hash(0x0564 0564 )  
hash(0x0645 0645 )  
hash(0x0726 0726 )  
hash(0x0807 0807 )
```

Equivalent Substring Attack on xquic

The nine 4-byte arrays all hash to the same value



```
hash(0x00ff 02c1 )  
hash(0x01e0 01e0 )  
hash(0x02c1 00ff )  
hash(0x03a2 0807 )  
hash(0x0483 0564 )  
hash(0x0564 0483 )  
hash(0x0645 03a2 )  
hash(0x0726 0645 )  
hash(0x0807 0726 )
```

Equivalent Substring Attack on xquic

```
00ff 0000 0000 0000 0000 0000  
01e0 0000 0000 0000 0000 0000  
02c1 0000 0000 0000 0000 0000  
03a2 0000 0000 0000 0000 0000  
0483 0000 0000 0000 0000 0000  
0564 0000 0000 0000 0000 0000  
0645 0000 0000 0000 0000 0000  
0726 0000 0000 0000 0000 0000  
0807 0000 0000 0000 0000 0000
```



9

Equivalent substring attack

9

colliding CIDs

Equivalent Substring Attack on xquic

00ff	0807	0000	0000	0000	0000	0000
01e0	00ff	0000	0000	0000	0000	0000
02c1	01e0	0000	0000	0000	0000	0000
03a2	02c1	0000	0000	0000	0000	0000
0483	03a2	0000	0000	0000	0000	0000
0564	0483	0000	0000	0000	0000	0000
0645	0564	0000	0000	0000	0000	0000
0726	0645	0000	0000	0000	0000	0000
0807	0726	0000	0000	0000	0000	0000



9 × 9

Equivalent substring attack

81

colliding CIDs

Equivalent Substring Attack on xquic

00ff	0807	0726	0000	0000	0000	0000
01e0	00ff	0807	0000	0000	0000	0000
02c1	01e0	00ff	0000	0000	0000	0000
03a2	02c1	01e0	0000	0000	0000	0000
0483	03a2	02c1	0000	0000	0000	0000
0564	0483	03a2	0000	0000	0000	0000
0645	0564	0483	0000	0000	0000	0000
0726	0645	0564	0000	0000	0000	0000
0807	0726	0645	0000	0000	0000	0000



9 × 9 × 9

Equivalent substring attack

729

colliding CIDs

QUIC CIDs MUST be 8-20 bytes long

Equivalent Substring Attack on xquic

00ff 0807 0726 0645 0564 0483
01e0 00ff 0807 0726 0645 0564
02c1 01e0 00ff 0807 0726 0645
03a2 02c1 01e0 00ff 0807 0726
0483 03a2 02c1 01e0 00ff 0807
0564 0483 03a2 02c1 01e0 00ff
0645 0564 0483 03a2 02c1 01e0
0726 0645 0564 0483 03a2 02c1
0807 0726 0645 0564 0483 03a2



9 × 9 × 9 × 9 × 9 × 9

Equivalent substring attack

531,441

colliding **12-byte** CIDs

botpaul@ubuntu ~ /D/h/x/cdemo> █

```
botpaul@ubuntu ~/D/h/x/cdemo> cat gen_collision.py
```

```
import itertools

# List of 2-byte hex strings
hex_values = ["00ff", "01e0", "02c1",
               "03a2", "0483", "0564",
               "0645", "0726", "0807"]

# Generate all 6-length permutations (with repetition)
for combo in itertools.product(hex_values, repeat=6):
    # Concatenate and print the result
    print("".join(combo))
botpaul@ubuntu ~/D/h/x/cdemo>
```

```
080708070807080707260807  
0807080708070807080700ff  
0807080708070807080701e0  
0807080708070807080702c1  
0807080708070807080703a2  
080708070807080708070483  
080708070807080708070564  
080708070807080708070645  
080708070807080708070726  
080708070807080708070807
```

```
botpaul@ubuntu ~ /D/h/x/cdemo> time python3 gen_collision.py > collisions.txt
```

```
Executed in 102.27 millis      fish          external  
  usr time  74.31 millis      0.01 millis    74.30 millis  
  sys time  25.23 millis     1.13 millis    24.10 millis
```

```
botpaul@ubuntu ~ /D/h/x/cdemo> wc -l collisions.txt
```

```
531441 collisions.txt
```

```
botpaul@ubuntu ~ /D/h/x/cdemo> █
```

Computing Collisions in xquic

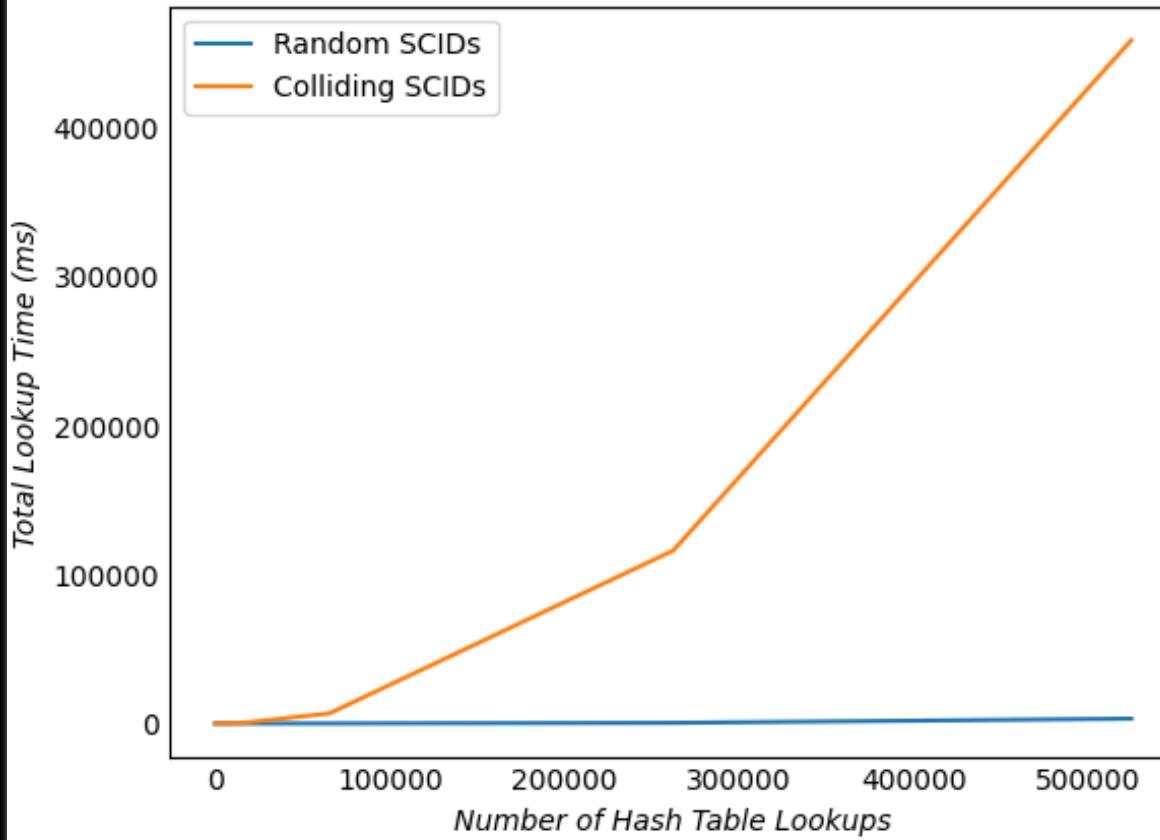
- Meet-in-the-middle attack
 - Build table of prefixes
 - Sample target hash
 - Draw random suffixes
 - Compute reverse intermediate hash
 - If match, input prefix+suffix is a hash collision
- SMT solver
- Linear property of the hash function



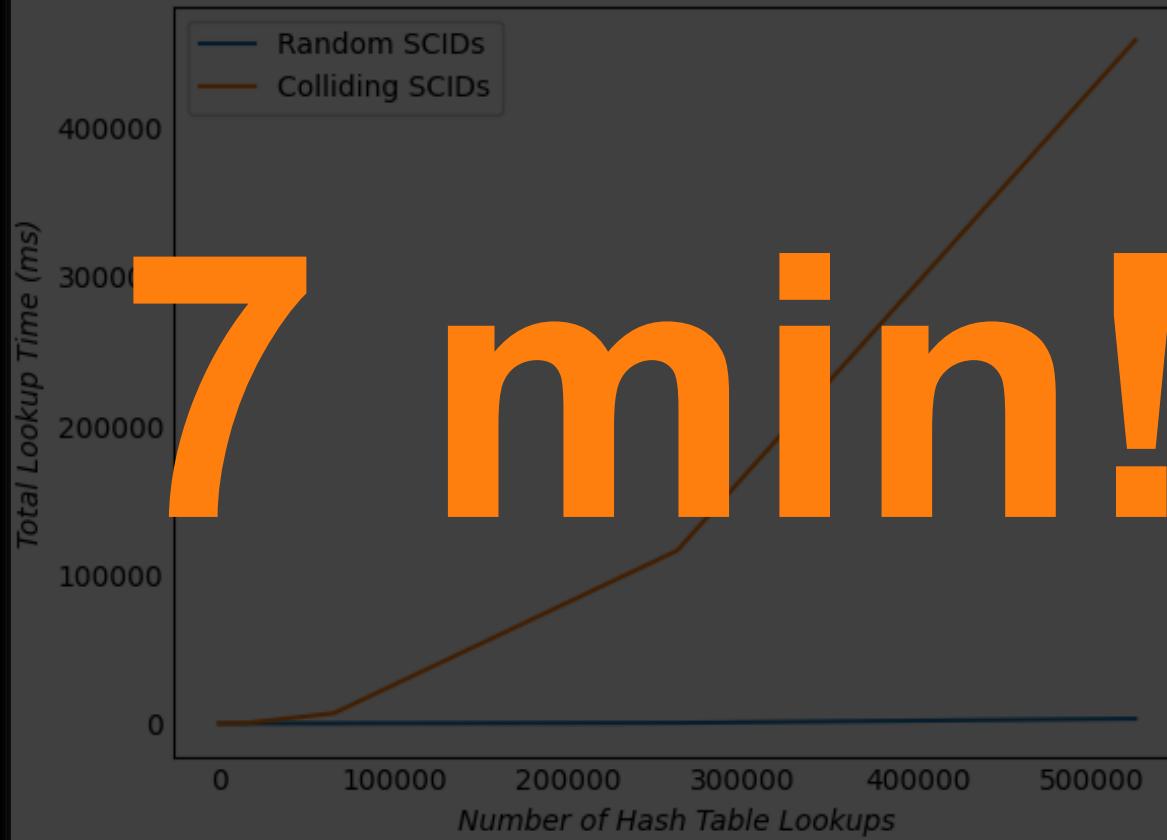
<https://github.com/pbottine/cut-to-the-quic>

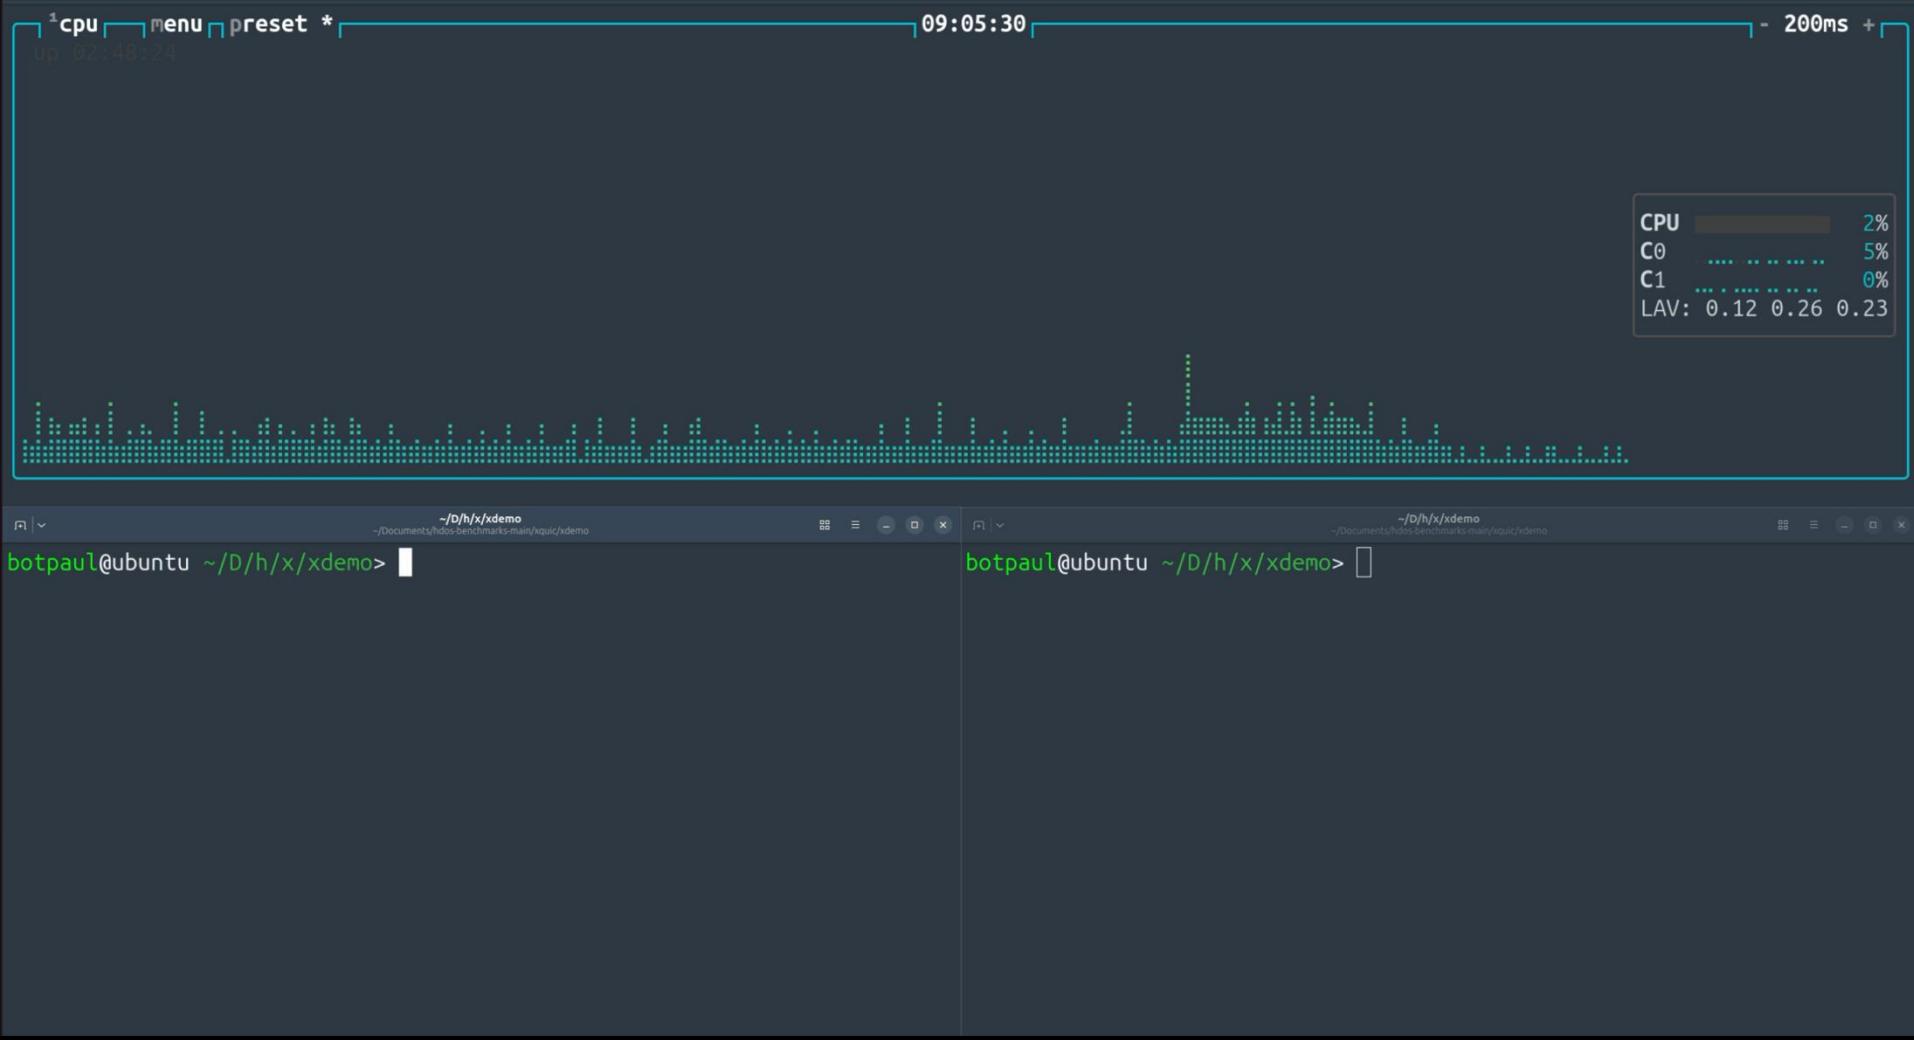
$$\text{hash} = c_n + c_{n-1} * 31 + c_{n-2} * 31^2 + \dots + c_1 * 31^{n-1}$$

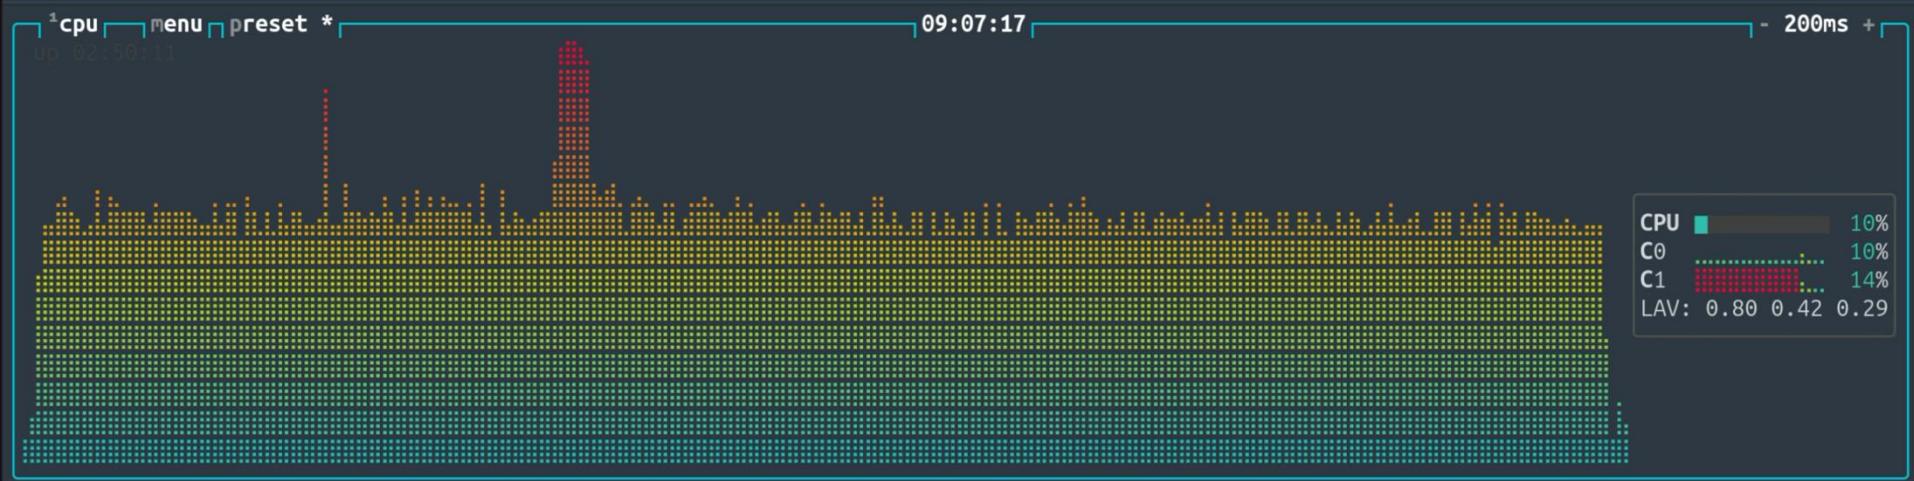
xquic - Cumulative Hash Table Lookup Time



xquic - Cumulative Hash Table Lookup Time







```
~/D/h/x/xdemo
~/Documents/hdos-benchmarks-main/xquic/xdemo
3154360FA0CA6A67C82BEEB5
73D9DCC721643A604D404CC7
botpaul@ubuntu ~/D/h/x/xdemo> time ./demo scid_random.txt 100000
```

```
Executed in 345.94 millis    fish      external
  usr time 339.44 millis    0.00 micros 339.44 millis
  sys time  6.41 millis   404.00 micros   6.01 millis
```

```
botpaul@ubuntu ~/D/h/x/xdemo> time ./demo scid_random.txt 200000
```

```
Executed in 1.11 secs    fish      external
  usr time 997.39 millis    0.00 micros 997.39 millis
  sys time  6.44 millis   447.00 micros   5.99 millis
```

```
botpaul@ubuntu ~/D/h/x/xdemo> █
```

```
~/D/h/x/xdemo
~/Documents/hdos-benchmarks-main/xquic/xdemo
00F800F900FA00FB00FC0805
00F800F900FA00FB01DD00FD
botpaul@ubuntu ~/D/h/x/xdemo> time ./demo scid_coll.txt 100000
```

```
Executed in 11.97 secs    fish      external
  usr time 11.97 secs    0.00 micros 11.97 secs
  sys time  0.00 secs   548.00 micros   0.00 secs
```

```
botpaul@ubuntu ~/D/h/x/xdemo> time ./demo scid_coll.txt 200000
```

```
Executed in 50.23 secs    fish      external
  usr time 49.98 secs    0.00 micros 49.98 secs
  sys time  0.01 secs   538.00 micros   0.01 secs
```

```
botpaul@ubuntu ~/D/h/x/xdemo> █
```

Hash DoS Ingredients



Hash DoS Ingredients



Attacker-controlled input

CVE-2025-
More weak hash functions?

XXH32 used in lsquic

xxHash - Extremely fast hash algorithm

xxHash is an Extremely fast Hash algorithm, processing at RAM speed limits. Code is highly portable, and produces hashes identical across all platforms (little / big endian). The library includes the following algorithms :

- XXH32 : generates 32-bit hashes, using 32-bit arithmetic
- XXH64 : generates 64-bit hashes, using 64-bit arithmetic
- XXH3 (since v0.8.0) : generates 64 or 128-bit hashes, using vectorized arithmetic. The 128-bit variant is called XXH128.

All variants successfully complete the [SMHasher](#) test suite which evaluates the quality of hash functions (collision, dispersion and randomness). Additional tests, which evaluate more thoroughly speed and collision properties of 64-bit hashes, [are also provided](#).

Source: <https://github.com/Cyan4973/xxHash>

```

class XXHash32 {
public:
    explicit XXHash32(uint32_t seed) {
        state[0] = seed + Prime1 + Prime2;
        state[1] = seed + Prime2;
        state[2] = seed;
        state[3] = seed - Prime1;
        bufferSize = 0;
        totalLength = 0;
    }
    bool add(const void* input, uint64_t length) {
        if (!input || length == 0)
            return false;
        totalLength += length;
        // SNIP
    }

private:
    static const uint32_t Prime1 = 2654435761U;
    static const uint32_t Prime2 = 2246822519U;
    static const uint32_t Prime3 = 3266489917U;
    static const uint32_t Prime4 = 668265263U;
    static const uint32_t Prime5 = 374761393U;

    static const uint32_t MaxBufferSize = 15+1;
    uint32_t state[4]; // state[2] == seed if totalLength < MaxBufferSize
    unsigned char buffer[MaxBufferSize];
    unsigned int bufferSize;
    uint64_t totalLength;

    static inline uint32_t rotateLeft(uint32_t x, unsigned char bits) {
        return (x << bits) | (x >> (32 - bits));
    }

    static inline void process(const void* data, uint32_t& state0,
        uint32_t& state1, uint32_t& state2, uint32_t& state3) {
        const uint32_t* block = (const uint32_t*) data;
        state0 = rotateLeft(state0 + block[0] * Prime2, 13) * Prime1;
        state1 = rotateLeft(state1 + block[1] * Prime2, 13) * Prime1;
        state2 = rotateLeft(state2 + block[2] * Prime2, 13) * Prime1;
        state3 = rotateLeft(state3 + block[3] * Prime2, 13) * Prime1;
    }
}

```

XXH32

```

uint32_t hash() const {
    uint32_t result = (uint32_t)totalLength;

    // fold 128 bit state into one single 32 bit value
    if (totalLength >= MaxBufferSize)
        result += rotateLeft(state[0], 1) +
                  rotateLeft(state[1], 7) +
                  rotateLeft(state[2], 12) +
                  rotateLeft(state[3], 18);
    else
        // internal state wasn't set in add(), therefore original seed is
        // still stored in state2
        result += state[2] + Prime5;

    // process remaining bytes in temporary buffer
    const unsigned char* data = buffer;

    // point beyond last byte
    const unsigned char* stop = data + bufferSize;

    // at least 4 bytes left ? => eat 4 bytes per step
    for (; data + 4 <= stop; data += 4)
        result = rotateLeft(result + *(uint32_t*)data * Prime3, 17) *
                 Prime4;

    // take care of remaining 0..3 bytes, eat 1 byte per step
    while (data != stop)
        result = rotateLeft(result + (*data++) * Prime5, 11) * Prime1;

    // mix bits
    result ^= result >> 15;
    result *= Prime2;
    result ^= result >> 13;
    result *= Prime3;
    result ^= result >> 16;
    return result;
}

```

```

class XXHash32 {
public:
    explicit XXHash32(uint32_t seed) {
        state[0] = seed + Prime1 + Prime2;
        state[1] = seed + Prime2;
        state[2] = seed;
        state[3] = seed - Prime1;
        bufferSize = 0;
        totalLength = 0;
    }

    bool add(const void* input, uint64_t length) {
        if (!input || length == 0)
            return false;
        totalLength += length;
        // SNIP
    }

private:
    static const uint32_t Prime1 = 2654435761U;
    static const uint32_t Prime2 = 2246822519U;
    static const uint32_t Prime3 = 3266489917U;
    static const uint32_t Prime4 = 668265263U;
    static const uint32_t Prime5 = 374761393U;

    static const uint32_t MaxBufferSize = 15+1;
    uint32_t state[4]; // state[2] == seed if totalLength < MaxBufferSize
    unsigned char buffer[MaxBufferSize];
    unsigned int bufferSize;
    uint64_t totalLength;

    static inline uint32_t rotateLeft(uint32_t x, unsigned char bits) {
        return (x << bits) | (x >> (32 - bits));
    }

    static inline void process(const void* data, uint32_t& state0,
        uint32_t& state1, uint32_t& state2, uint32_t& state3) {
        const uint32_t* block = (const uint32_t*) data;
        state0 = rotateLeft(state0 + block[0] * Prime2, 13) * Prime1;
        state1 = rotateLeft(state1 + block[1] * Prime2, 13) * Prime1;
        state2 = rotateLeft(state2 + block[2] * Prime2, 13) * Prime1;
        state3 = rotateLeft(state3 + block[3] * Prime2, 13) * Prime1;
    }
}

```

XXH32

```

uint32_t hash() const {
    uint32_t result = (uint32_t)totalLength;

    // fold 128 bit state into one single 32 bit value
    if (totalLength >= MaxBufferSize)
        result += rotateLeft(state[0], 1) +
                  rotateLeft(state[1], 7) +
                  rotateLeft(state[2], 12) +
                  rotateLeft(state[3], 18);
    else
        // internal state wasn't set in add(), therefore original seed is
        // still stored in state2
        result += state[2] + Prime5;

    // process remaining bytes in temporary buffer
    const unsigned char* data = buffer;

    // point beyond last byte
    const unsigned char* stop = data + bufferSize;

    // at least 4 bytes left ? => eat 4 bytes per step
    for (; data + 4 <= stop; data += 4)
        result = rotateLeft(result + *(uint32_t*)data * Prime3, 17) *
                 Prime4;

    // take care of remaining 0..3 bytes, eat 1 byte per step
    while (data != stop)
        result = rotateLeft(result + (*data++) * Prime5, 11) * Prime1;

    // mix bits
    result ^= result >> 15;
    result *= Prime2;
    result ^= result >> 13;
    result *= Prime3;
    result ^= result >> 16;
    return result;
}

```

Simplified XXH32

```
uint32_t hash() const {
    uint32_t result = (uint32_t)totalLength;

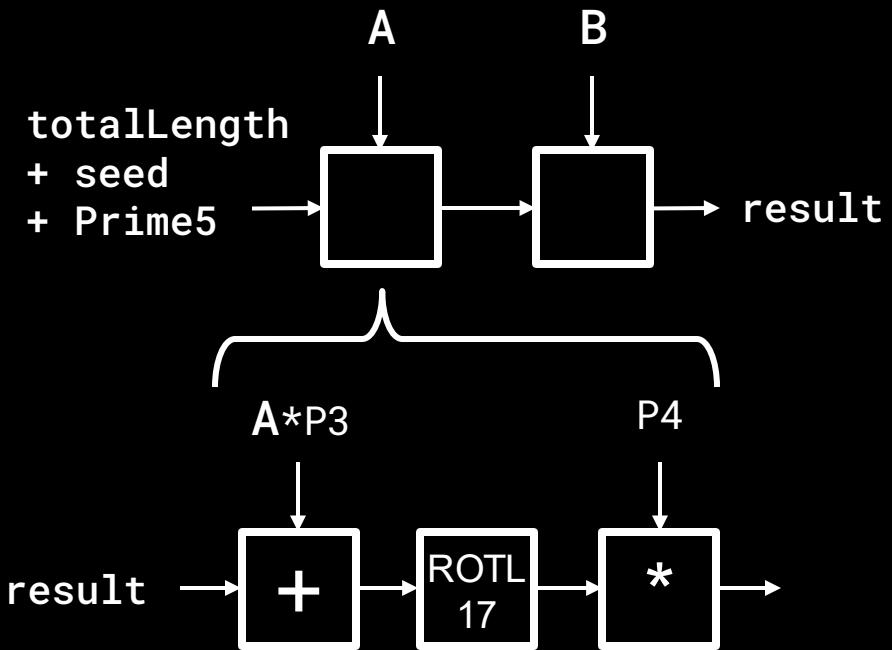
    result += seed + Prime5;

    // process remaining bytes in temporary buffer
    const unsigned char* data = buffer;

    // point beyond last byte
    const unsigned char* stop = data + bufferSize;

    // at least 4 bytes left ? => eat 4 bytes per step
    for (; data + 4 <= stop; data += 4) {
        result += *(uint32_t*)data * Prime3;
        result = rotateLeft(result, 17);
        result *= Prime4;
    }
    return result;
}
```

Simplified XXH32



```
uint32_t hash() const {
    uint32_t result = (uint32_t)totalLength;

    result += seed + Prime5;

    // process remaining bytes in temporary buffer
    const unsigned char* data = buffer;

    // point beyond last byte
    const unsigned char* stop = data + bufferSize;

    // at least 4 bytes left ? => eat 4 bytes per step
    for (; data + 4 <= stop; data += 4) {
        result += *(uint32_t*)data * Prime3;
        result = rotateLeft(result, 17);
        result *= Prime4;
    }
    return result;
}
```

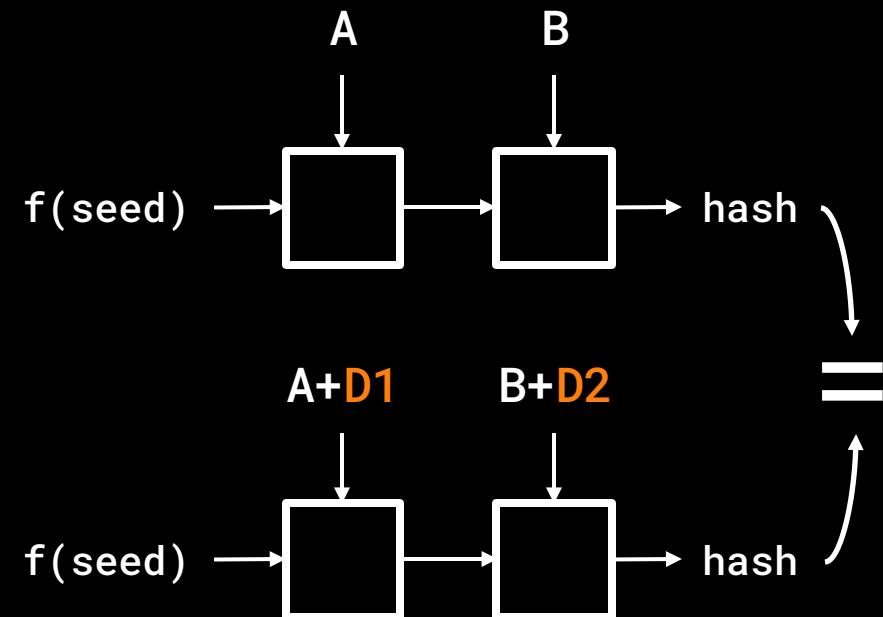
Simplified hash assuming `len(input) = 8`

Differential Cryptanalysis of XXH32

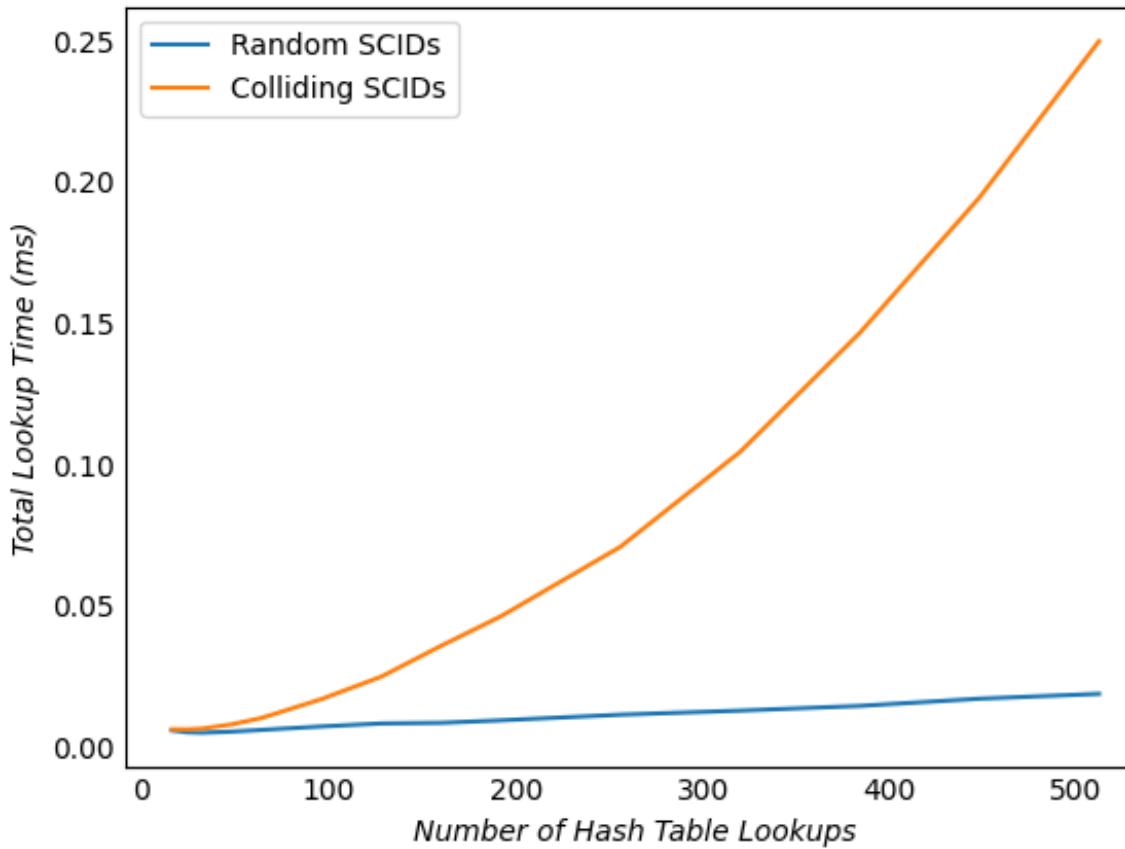
- Given input $A||B$, we can find pairs of deltas $(D1, D2)$ such that

$$\text{XXH32}(A||B) = \text{XXH32}(A+D1||B+D2)$$

- Independent of the seed; collisions work for any seed!
- Exhaustive search yields over 100k of $(D1, D2)$ pairs per input $A||B$

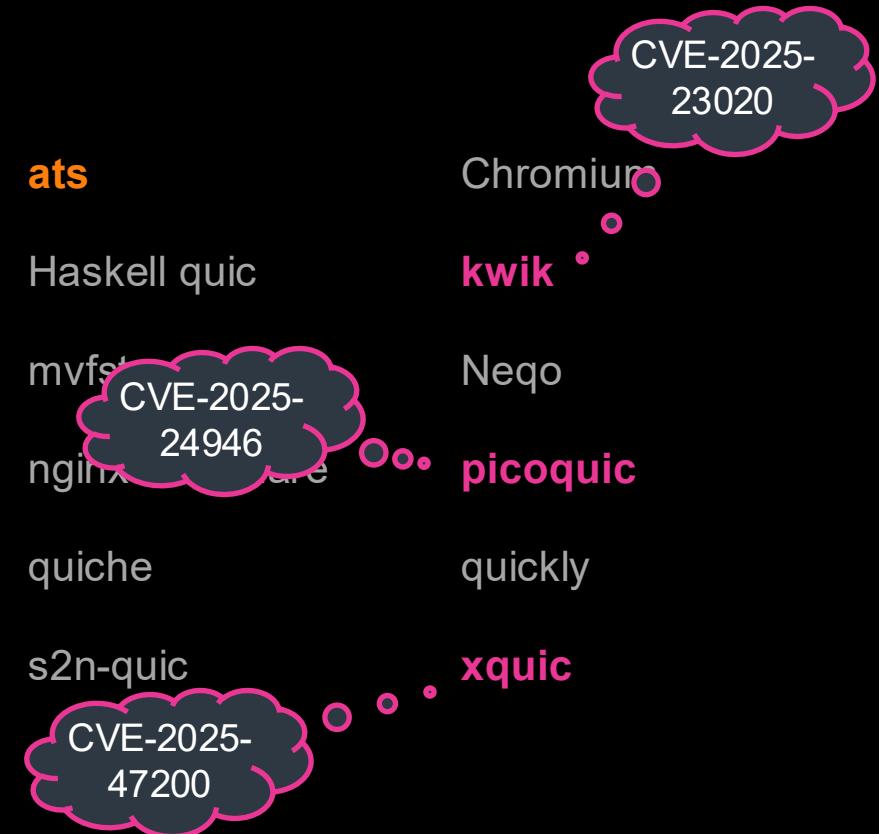


Isquic - Cumulative Hash Table Lookup Time

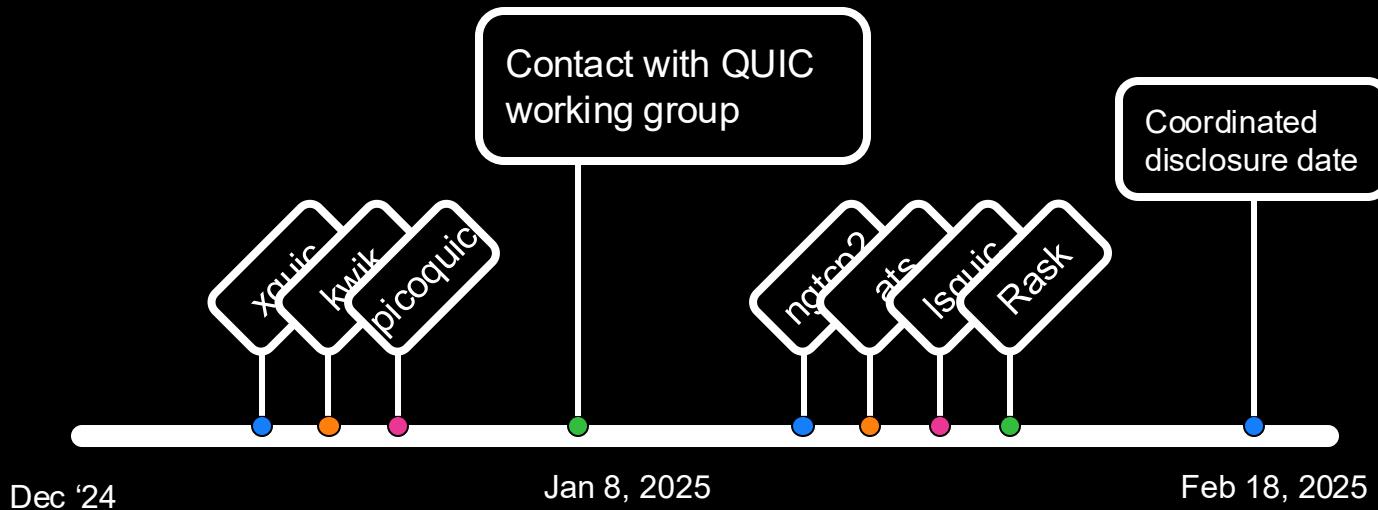


Results

aioquic	AppleQUIC	ats	Chromium
f5	Cloudflare QUIC (f5)	Haskell quic	kwik
lsquic	MsQuic	mvfst	Neqo
ngtcp2	nginx	nginx	picoquic
Pluginized QUIC	quant	quiche	quickly
Quinn	quic-go	s2n-quic	xquic



Disclosure Timeline



Interesting Notes

- Ericsson Rask, an experimental, pre-GA MP-QUIC implementation used Rust
- Chromium had noticed and fixed the issue in 2019
- New CVE was assigned after disclosure to Netty QUIC



CVE-2025-29908 Detail

AWAITING ANALYSIS

This CVE record has been marked for NVD enrichment efforts.

Description

Netty QUIC codec is a QUIC codec for netty which makes use of quiche. An issue was discovered in the codec. A hash collision vulnerability (in the hash map used to manage connections) allows remote attackers to cause a considerable CPU load on the server (a Hash DoS attack) by initiating connections with colliding Source Connection IDs (SCIDs). This vulnerability is fixed in 0.0.71.Final.

Black Hat Sound Bytes

Attacker-controlled input

Protocol Designers:

- Design protocols with care
- Prevent attacker-controlled input
- Consider RFC 9414, 9415 and 9416



Vulnerability researchers:

- Keep findings Hash DoS attacks!

Weak hash function

Protocol Implementers:

- Use languages with built-in Hash DoS protections
- Use hash functions that provide security
- Don't roll your own

Thank you

- Trail of Bits (my employer)
- Javed Samuel @ NCC Group
- Lucas Pardue for coordinating disclosure process
- BlackHat for having me
- Phil Young for coaching



<https://github.com/pbottine/cut-to-the-quic>

Contact info

Name: Paul Bottinelli

Email: paul.bottinelli@trailofbits.com

LinkedIn: paulbottinelli

Bsky: @botpaul

