

Buckle Up, Buttercup: Our Experience Competing in the AI Cyber Challenge

9 August 2025

Our Team



Michael D. Brown

Overall Team Lead
Lead Designer of Buttercup



Ian Smith

Vuln Discovery Lead
Co-Designer of Buttercup



Ronald Eytchison

AI-Based Seed Generation Lead

Our Team



Henrik Brodin

Orchestration Lead
(Finals)



Eric Kilmer

Orchestration Co-Lead
(Semi-Finals)



Francesco Bertolaccini

Orchestration Co-Lead
(Semi-Finals)

Our Team



Riccardo Schirrone

Patcher Lead



Evan Downing

Contextualization Lead



Boyan Milanov

System Developer

Our Team



Alessandro Gario

Challenge Creator
(Internal Red Team)



Brad Swain

Challenge Creator
(Internal Red Team)

Our Team

Akshay Kumar

Challenge Creation
(Semi-finals)

Will Tan

Systems Developer
(Semi-finals)


Alan Cao

Systems Developer
(Semi-finals)


A Brief Origin Story

AI Cyber Challenge (AIxCC)


AIxCC is a competition to design a novel automated AI system (CRS) that can find and patch bugs in real-world open-source software.




AIxCC Competition Structure



AIxCC Competition Structure



AIxCC Competition Structure



Buttercup's Design

Our Approach

Guiding Principles


- Conventional software analysis works really well for certain problems.
- AI/ML-based analysis works really well for certain problems.
- Often, one approach works well where the other does not.

Break the problem down, use the best technique to solve each sub-problem.
Don't expect LLMs to do things they aren't good at!


Problem Breakdown

- 1) Discover / prove existence of vulnerabilities
- 2) Contextualize vulnerabilities
- 3) Create and Validate patches
- 4) Orchestrate these tasks to:
 - a) Effectively allocate resources
 - b) Maximize score


CRS Architecture (Concept Paper)



CRS Architecture



CRS Architecture (Competition)



Buttercup in the Semifinals

ACHIEVEMENTS UNLOCKED

 TEAM NAME: Trail of Bits	 TEAM NAME: 42-b3yond-6ug	 TEAM NAME: Theori	 TEAM NAME: Theori	 TEAM NAME: 42-b3yond-6ug	 TEAM NAME: 42-b3yond-6ug	 TEAM NAME: 42-b3yond-6ug	 TEAM NAME: Theori	 TEAM NAME: Theori	 TEAM NAME: Trail of Bits	 TEAM NAME: Trail of Bits	
 TEAM NAME: Team Atlanta	 TEAM NAME: Theori	 TEAM NAME: Trail of Bits	 TEAM NAME: 42-b3yond-6ug	 TEAM NAME: Trail of Bits	 TEAM NAME: Theori	 TEAM NAME: Theori	 TEAM NAME: Theori	 TEAM NAME: Trail of Bits	 TEAM NAME: Theori	 TEAM NAME: Theori	
 TEAM NAME: Trail of Bits	 TEAM NAME: Shellphish	 TEAM NAME: 42-b3yond-6ug	 TEAM NAME: all-you-need-is-a-f...	 TEAM NAME: Trail of Bits	 TEAM NAME: 42-b3yond-6ug	 TEAM NAME: Trail of Bits	 TEAM NAME: 42-b3yond-6ug	 TEAM NAME: Team Atlanta	 TEAM NAME: all-you-need-is-a-f...	 TEAM NAME: Theori	
 TEAM NAME: 42-b3yond-6ug	 TEAM NAME: Team Atlanta	 TEAM NAME: 42-b3yond-6ug	 TEAM NAME: Theori	 TEAM NAME: Theori	 TEAM NAME: Trail of Bits	 TEAM NAME: Theori	 TEAM NAME: Team Atlanta	 TEAM NAME: Team Atlanta	 TEAM NAME: Lacrosse	 TEAM NAME: Theori	
 TEAM NAME: Team Atlanta		 TEAM NAME: Theori									



Tika

Status: Vulnerable

Jenkins

Status: Vulnerable

Linux Kernel

Status: Vulnerable

Sqlite3

Status: Vulnerable




Vulnerable



Discovery




Patched



Nginx
Status: Vulnerable



Performance by CWE type



Buttercup 2.0

How did Buttercup evolve for the finals?

Lessons Learned from semi-finals:

- Validated our overall approach
- Need better testing / handling of Java challenges
- CWE-type specific seed-generation may have helped

Rule changes for finals:

- Massive scale and budget (time, compute, and AI) increases
- Several exhibition rounds
- More complex scoring (SARIFs, bundles, duplication penalties)
- Custom AI/ML models allowed

Building Buttercup 2.0


Buttercup 2.0 is essentially a from-scratch rebuild.

Driven by need for:

- more technically complex analysis components
- ability to easily change scale / cost of deployment for various rounds
- high degree of reliability and robustness to errors

Still, our high-level Buttercup remained the same as the semi-finals

CRS Architecture (Competition)



Buttercup 2.0 Technical Details

Orchestration – Submission Processing

Filter

Group by stacktrace

Group by patch

Monitor

Vulnerability discovery produces many PoVs - filter stack traces already seen

Group PoVs with similar stack traces - examples of the same underlying vulnerability.

Group PoVs remediated by the same patch - same underlying vulnerability

As new PoVs come in merge by fuzzy stack match and patches. Rebuild bundles as needed.

PoV - Proof of Vulnerability

Vulnerability Discovery

- **Strategy: Combine fuzzing and LLM input generation**
- Use standard OSS-Fuzz fuzzers:
 - LibFuzzer for C/C++
 - Jazzer for Java
- Fuzzer bots sample active harnesses to run short fuzz campaigns
- Fuzzing corpus:
 - Merger bots merge a fuzzer bot's local corpus to the shared corpus
 - LLM input generation also submits to the corpus

Vulnerability Discovery: LLM “seed-gen”

Design

- Several tasks that use LLMs to create seeds and/or PoVs
- All tasks use tools to collect context from the codebase before generating inputs


Goal 1: Support Fuzzing

- **Init task:** Bootstrap fuzzer with initial seed inputs that exercise harness
- **Explore task:** Increase coverage for a target function

Goal 2: Independently Find Bugs

- **Vuln discovery task:** Identify and validate vulnerabilities in target to create PoVs
 - Most expensive task to thoroughly explore code and test hypotheses

Contextualization




- Constructs program model using CodeQuery + Tree-sitter
- Supports querying program properties (functions & types)
- Called by LLMs from **Seed Generator** and **Patcher** using LangGraph's Tool library


Patcher

- LLM-based multi-agent system
 - Software, Security, and Quality Engineer Agents working together
- Programmatic agents hand-off
 - Data flow between agent is (mostly) deterministic
 - More control over the process
 - Error handling relies on LLMs to determine resolution steps
- **Implementation**
 - Less than 6K LOC, Python
 - LangChain/LangGraph
 - Preferred model: OpenAI/GPT-4.1

Patcher: flow



Patcher: patch creation



Buttercup in the Finals

How did Buttercup do in Exhibition Rounds?

Buttercup was the best performing CRS in Round 1:

- Found and patched a vulnerability in both challenges with 100% accuracy
- Used only ~\$1000 of available \$30,000 budget

But we crashed hard in Round 2:

- Issue with filename length in vulnerability discovery component
- Caused a hard failure after only 3/18 challenges were processed
- We later reproduced Round 2 and Buttercup was successful on all challenges

And bounced back in Round 3:

- Buttercup found and/or patched vulnerabilities in 20/26 challenges!

How did Buttercup do in the scored round?

Buttercup came in second place, winning \$3 million!

- Found 28 vulnerabilities, patched 19
- Used only ~\$40,000 of available budget
- ~90% Accuracy
- Found at least one PoV no one else did
- Found at least one non-synthetic vulnerability

Keys to success:

- Accuracy
- Scoring well across all tasks

I want to try Buttercup!

You're in Luck....

Buttercup is Open Source!

The exact code we submitted for the semi-finals and finals code is available on our company github organization!

- Buttercup 1.0 <https://github.com/trailofbits/asc-buttercup>
- Buttercup 2.0 <https://github.com/trailofbits/afc-buttercup>

Fair warning: Buttercup was designed to run on competition infrastructure and at massive scale, so this version of Buttercup isn't terribly user friendly...

And we'll do you one better!

A standalone variant of Buttercup is also available!

We've also created a version of Buttercup that runs on commodity (laptop) and typical server-grade hardware. You can check it out at:

- Buttercup standalone <https://github.com/trailofbits/buttercup>

Enjoy!

Thanks for Coming!