



Solang Parser and Semantic Analysis

Security Assessment (Summary Report)

September 6, 2023

Prepared for:

Sean Young

Solana Labs

Prepared by: **Anders Helsing and Samuel Moelius**

About Trail of Bits

Founded in 2012 and headquartered in New York, Trail of Bits provides technical security assessment and advisory services to some of the world's most targeted organizations. We combine high-end security research with a real-world attacker mentality to reduce risk and fortify code. With 100+ employees around the globe, we've helped secure critical software elements that support billions of end users, including Kubernetes and the Linux kernel.

We maintain an exhaustive list of publications at <https://github.com/trailofbits/publications>, with links to papers, presentations, public audit reports, and podcast appearances.

In recent years, Trail of Bits consultants have showcased cutting-edge research through presentations at CanSecWest, HCSS, Devcon, Empire Hacking, GrrCon, LangSec, NorthSec, the O'Reilly Security Conference, PyCon, REcon, Security BSides, and SummerCon.

We specialize in software testing and code review projects, supporting client organizations in the technology, defense, and finance industries, as well as government entities. Notable clients include HashiCorp, Google, Microsoft, Western Digital, and Zoom.

Trail of Bits also operates a center of excellence with regard to blockchain security. Notable projects include audits of Algorand, Bitcoin SV, Chainlink, Compound, Ethereum 2.0, MakerDAO, Matic, Uniswap, Web3, and Zcash.

To keep up to date with our latest news and announcements, please follow [@trailofbits](#) on Twitter and explore our public repositories at <https://github.com/trailofbits>. To engage us directly, visit our "Contact" page at <https://www.trailofbits.com/contact>, or email us at info@trailofbits.com.

Trail of Bits, Inc.

228 Park Ave S #80688

New York, NY 10003

<https://www.trailofbits.com>

info@trailofbits.com

Notices and Remarks

Copyright and Distribution

© 2023 by Trail of Bits, Inc.

All rights reserved. Trail of Bits hereby asserts its right to be identified as the creator of this report in the United Kingdom.

This report is considered by Trail of Bits to be public information; it is licensed to Solana Labs under the terms of the project statement of work and has been made public at Solana Labs' request. Material within this report may not be reproduced or distributed in part or in whole without the express written permission of Trail of Bits.

The sole canonical source for Trail of Bits publications is the [Trail of Bits Publications page](#). Reports accessed through any source other than that page may have been modified and should not be considered authentic.

Test Coverage Disclaimer

All activities undertaken by Trail of Bits in association with this project were performed in accordance with a statement of work and agreed upon project plan.

Security assessment projects are time-boxed and often reliant on information that may be provided by a client, its affiliates, or its partners. As a result, the findings documented in this report should not be considered a comprehensive list of security issues, flaws, or defects in the target system or codebase.

Trail of Bits uses automated testing techniques to rapidly test the controls and security properties of software. These techniques augment our manual security review work, but each has its limitations: for example, a tool may not generate a random edge case that violates a property or may not fully complete its analysis during the allotted time. Their use is also limited by the time and resource constraints of a project.

Table of Contents

About Trail of Bits	1
Notices and Remarks	2
Table of Contents	3
Executive Summary	4
Project Summary	6
Project Goals	7
Project Targets	8
Project Coverage	9
Summary of Findings	11
Detailed Findings	12
1. Reliance on deprecated/unmaintained dependencies	12
2. Reliance on outdated dependencies	14
3. Insufficient linter use	16
4. Over-reliance on casts that could fail	18
5. Rational evaluation code always returns errors	20
6. Code duplication	22
7. Risk of arithmetic underflow in lexer	25
8. Excessive use of allow for lalrpop-generated code	27
9. No explicit tests for operator precedence	28
A. Vulnerability Categories	29
B. Non-Security-Related Findings	32

Executive Summary

Engagement Overview

Solana Labs engaged Trail of Bits to review the security of the Solang compiler's parsing and semantic analysis phases. From March 13 to March 17, 2023, a team of two consultants conducted a security review of the client-provided source code, with two person-weeks of effort. Details of the project's timeline, test targets, and coverage are provided in subsequent sections of this report.

Project Scope

Our testing efforts were focused on the identification of flaws that could result in a compromise of confidentiality, integrity, or availability of the target system. We conducted this audit with full knowledge of the target system, including access to the source code and documentation. We performed static and dynamic testing of the target system and its codebase, using both automated and manual processes.

Summary of Findings

The audit uncovered only issues of informational or undetermined severity that could impact system confidentiality, integrity, or availability. A summary of the findings is provided on the next page.

Summary of Recommendations

Trail of Bits recommends addressing all of the findings in this report, but implementing the following recommendations would have the largest impact on the codebase's security:

- Reduce the amount of code duplication. There appears to be a considerable amount of code duplication within this codebase. Duplicate code can result in incomplete fixes or inconsistent behavior (e.g., if the code is modified in one location but not in all). (TOB-SOLANG-6)
- Increase the use of Clippy. Currently, Clippy appears to be run with only its default set of lints enabled (`clippy :all`). However, several pedantic lints appear to produce valid warnings when run on the code. (TOB-SOLANG-3)
- Review all casts and eliminate those that are unnecessary. There are at least 63 casts within the source code, which seems excessive. Consider using `#[deny(clippy::as_conversions)]` and selectively allowing the lint where casts are necessary and known to be safe. (TOB-SOLANG-4)
- Incorporate fuzzing into the software development process. Nearly all of the data a compiler operates on is user-controlled. For this reason, compilers often benefit considerably from fuzzing. (TOB-SOLANG-7)

EXPOSURE ANALYSIS

<i>Severity</i>	<i>Count</i>
High	0
Medium	0
Low	1
Informational	7
Undetermined	1

CATEGORY BREAKDOWN

<i>Category</i>	<i>Count</i>
Data Validation	2
Error Reporting	1
Patching	3
Testing	3

Project Summary

Contact Information

The following manager was associated with this project:

Jeff Braswell, Project Manager
jeff.braswell@trailofbits.com

The following engineering director was associated with this project:

Josselin Feist, Engineering Director, Blockchain
josselin.feist@trailofbits.com

The following engineers were associated with this project:

Anders Helsing , Consultant anders.helsing@trailofbits.com	Samuel Moelius , Consultant samuel.moelius@trailofbits.com
--	--

Project Timeline

The significant events and milestones of the project are listed below.

Date	Event
March 13, 2023	Pre-project kickoff call
March 20, 2023	Delivery of report draft
March 20, 2023	Report readout meeting
September 6, 2023	Delivery of final report

Project Goals

The engagement was scoped to provide a security assessment of the Solang compiler's parsing and semantic analysis phases. Specifically, we sought to answer the following non-exhaustive list of questions:

- Are there unintended parser differentials between the Solang and Solidity compilers?
- Does the parser/lexer interpret Solidity files correctly?
- Are operations involving types implemented correctly?
- Are the existing testing strategies sufficient, or could they be extended?

Project Targets

The engagement involved a review and testing of the following targets.

Solang Parser

Repository	https://github.com/hyperledger/solang/tree/main/solang-parser
Version	52879197d1cbc0e83c1b44de70639256d73de105
Type	Rust
Platform	Solana

Solang Semantics

Repository	https://github.com/hyperledger/solang/tree/main/src/sema
Version	52879197d1cbc0e83c1b44de70639256d73de105
Type	Rust
Platform	Solana

Project Coverage

This section provides an overview of the analysis coverage of the review, as determined by our high-level engagement goals. Our approaches included the following:

- **Dependency review:** We ran `cargo upgrade --incompatible` over the codebase and reviewed the changes that the command applied to the `solana-ledger` crate.
- **Static analysis:** We ran static analysis tools and linters over the codebase and triaged the results.
- **Code review:** We reviewed the `lalrpop` definitions and the lexer from the `solang-parser` crate. We also began reviewing the semantic definitions, focusing on identifying problems stemming from type conversions and overflow conditions, but we were unable to complete this review.
- **Fuzzing:** We changed instances of silent integer truncation so that they panic on invalid conversions, such as expressions in the form `x as u8` into `x.try_into()`, and compiled the contracts in the `fuzzy-sol` repository. We also fuzzed the lexer.

Coverage Limitations

Because of the time-boxed nature of testing work, it is common to encounter coverage limitations. During this project, we were unable to perform comprehensive testing of the following system elements, which may warrant further review:

1. The semantic definitions (completion of our review for problems stemming from type conversions and overflow conditions)
2. Handling of data types that could differ between `solc` and `Solang`, including addresses, messages, and blocks
3. ABI encoding (`encodeWithSelector`, etc.)

We think that coverage of the above system elements could be completed with four person-weeks of effort. Note that this would not include the review of YUL.

We recommend that each of the following areas be the subject of their own audit, independent of the above.

1. **YUL:** As the YUL implementation becomes more complete, it should be audited to identify any security issues resulting from the use of inline Solidity assembly operating in a Solana context.

2. **Code generation:** The compiler's generation of the code should be audited to ensure that the Solidity contract code is correctly interpreted and that the generated Solana binary does not exhibit problems similar to those outlined in the [sealevel attacks](#).

Summary of Findings

The table below summarizes the findings of the review, including type and severity details.

ID	Title	Type	Severity
1	Reliance on deprecated/unmaintained dependencies	Patching	Informational
2	Reliance on outdated dependencies	Patching	Informational
3	Insufficient linter use	Testing	Informational
4	Over-reliance on casts that could fail	Data Validation	Undetermined
5	Rational evaluation code always returns errors	Error Reporting	Informational
6	Code duplication	Patching	Informational
7	Risk of arithmetic underflow in lexer	Data Validation	Low
8	Excessive use of allow for lalrpop-generated code	Testing	Informational
9	No explicit tests for operator precedence	Testing	Informational

Detailed Findings

1. Reliance on deprecated/unmaintained dependencies

Severity: Informational

Difficulty: High

Type: Patching

Finding ID: TOB-SOLANG-1

Target: Cargo.lock

Description

The Solang compiler relies on the `parity-wasm` crate, which has been declared “deprecated” by its author. Deprecated software is less likely to receive security updates. Hence, alternatives should be used.

```
Crate:    parity-wasm
Version:  0.42.2
Warning:  unmaintained
Title:    Crate `parity-wasm` deprecated by the author
Date:     2022-10-01
ID:       RUSTSEC-2022-0061
URL:      https://rustsec.org/advisories/RUSTSEC-2022-0061
Dependency tree:
parity-wasm 0.42.2
├─ wasmi-validation 0.4.1
│   └─ wasmi 0.11.0
│       └─ solang 0.2.2
└─ wasmi 0.11.0

Crate:    parity-wasm
Version:  0.45.0
Warning:  unmaintained
Title:    Crate `parity-wasm` deprecated by the author
Date:     2022-10-01
ID:       RUSTSEC-2022-0061
URL:      https://rustsec.org/advisories/RUSTSEC-2022-0061
Dependency tree:
parity-wasm 0.45.0
└─ solang 0.2.2
```

Figure 1.1: Results of running `cargo-audit`

Additionally, the Solang compiler relies on code generated by the `lalrpop` crate. In 2018, the crate’s author stated that he did not “have the time to devote to LALRPOP that it really deserves,” and sought to “**form a core team.**” However, it is unclear whether such a team

was formed. See, for example, issue #712 ("[Fork and add maintainers?](#)") from March 14, 2023.

However, at present, we do not know of a good alternative to recommend for `lalrpop`.

Exploit Scenario

Eve learns of a vulnerability in the code generated by `lalrpop`. Knowing that the Solang compiler relies on code generated by `lalrpop`, Eve exploits the compiler's users.

Recommendations

Short term, replace the use of `parity-wasm` with `wasm-tools`, which is recommended in the associated [RUSTSEC advisory](#). Additionally, keep an eye out for a fork of `lalrpop` or the announcement of a suitable alternative. Software that is actively maintained is more likely to receive security updates.

Long term, regularly run `cargo-audit` over the codebase. Doing so will help to identify vulnerable or unmaintained dependencies.

References

- [lalrpop issue #290](#), "form a core team"
- [lalrpop issue #712](#), "Fork and add maintainers?"
- [RUSTSEC-2022-0061](#)

2. Reliance on outdated dependencies

Severity: Informational

Difficulty: High

Type: Patching

Finding ID: TOB-SOLANG-2

Target: Cargo.toml

Description

Updated versions of many of the Solang compiler's dependencies are available. Because silent bug fixes are common, all dependencies should be periodically reviewed and updated wherever possible.

Note that some of these outdated dependencies have updated versions that are considered incompatible by Cargo; because of this, simply running `cargo update` will not cause them to be updated in the project's `Cargo.lock` file. Dependencies for which incompatible upgrades are available appear in table 2.1.

Dependency	Version currently in use	Latest version available
contract-metadata	1.5.1	2.1.0
tower-lsp	0.18	0.19.0
anchor-syn	0.26	0.27.0
wasmi	0.11	0.28.0
rand (rand_07)	0.7	0.8.5

Table 2.1: Dependencies for which incompatible upgrades are available

Exploit Scenario

Eve learns of a vulnerability in an outdated version of a Solang parser dependency. Knowing that Solana still relies on this outdated version, Eve exploits the vulnerability.

Recommendations

Short term, update the dependencies to their latest versions wherever possible. Verify that all unit tests pass following such updates. Document any reasons for not updating a dependency. Using out-of-date dependencies could mean critical bug fixes are missed.

Long term, regularly run `cargo upgrade --incompatible`. This will help ensure that the project stays up to date with its dependencies.

3. Insufficient linter use

Severity: Informational

Difficulty: High

Type: Testing

Finding ID: TOB-SOLANG-3

Target: Various source files

Description

The Solang parser appears to run Clippy with only the default set of lints enabled (`clippy:::all`). The maintainers should consider running additional lints, as many of them trigger warnings when enabled for the project.

Running Clippy with `-W clippy::pedantic` produces several hundred warnings, indicating that enabling additional lints could have a positive impact on the correctness of the codebase. Examples of such warnings appear in figures 3.1 through 3.4.

```
warning: used `cloned` where `copied` could be used instead
--> src/sema/contracts.rs:384:55
    |
384 | ...                override_specified.iter().cloned().collect();
    |                                           ^^^^^^^ help: try: `copied`
    |
= help: for further information visit
https://rust-lang.github.io/rust-clippy/master/index.html#cloned\_instead\_of\_copied
```

Figure 3.1: A sample warning produced by the `clippy::cloned_instead_of_copied` lint

```
warning: you seem to be using a `LinkedList`! Perhaps you meant some other data
structure?
--> src/sema/symtable.rs:85:12
    |
85 |     names: LinkedList<VarScope>,
    |             ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
    |
= help: a `VecDeque` might work
= help: for further information visit
https://rust-lang.github.io/rust-clippy/master/index.html#linkedlist
```

Figure 3.2: A sample warning produced by the `clippy::linkedlist` lint

```
warning: this argument is passed by value, but not consumed in the function body
--> src/sema/expression/function_call.rs:238:19
    |
```

```

238 |         function_nos: Vec<usize>,
    |                        ^^^^^^^^^^^^^ help: consider changing the type to: `&[usize]`
    |
    = help: for further information visit
https://rust-lang.github.io/rust-clippy/master/index.html#needless\_pass\_by\_value

```

Figure 3.3: A sample warning produced by the `clippy::needless_pass_by_value` lint

```

warning: this argument (1 byte) is passed by reference, but would be more efficient
if passed by value (limit: 8 byte)
--> src/sema/yul/expression.rs:100:12
100 |         value: &bool,
    |                ^^^^^ help: consider passing by value instead: `bool`
    |
    = help: for further information visit
https://rust-lang.github.io/rust-clippy/master/index.html#trivially\_copy\_pass\_by\_ref

```

Figure 3.4: A sample warning produced by the `clippy::trivially_copy_pass_by_ref` lint

Exploit Scenario

Eve uncovers a bug in the Solang parser. The bug might have been caught by the Solang compiler developers had additional lints been enabled.

Recommendations

Short term, review all of the warnings currently generated by Clippy's pedantic lints. Address those in figures 3.1 through 3.4, and any others for which it makes sense to do so. Taking these steps will produce cleaner code, which in turn will reduce the likelihood that the code contains bugs. (See also finding [TOB-SOLANG-4](#).)

Long term, take the following steps:

- Consider running Clippy with `-W clippy::pedantic` regularly. As demonstrated by the warnings in figures 3.1 through 3.4, the pedantic lints provide valuable suggestions.
- Regularly review Clippy lints that have been allowed to see whether they should still be given such an exemption. Allowing a Clippy lint unnecessarily could cause bugs to be missed.

4. Over-reliance on casts that could fail

Severity: **Undetermined**

Difficulty: **Low**

Type: Data Validation

Finding ID: TOB-SOLANG-4

Target: sema/expression/literals.rs

Description

The `cast_sign_loss` and `cast_possible_wrap` Clippy lints warn about casts that could fail. When run on the Solana parser, the lints produce 63 warnings. Moreover, user-controlled data can cause at least one of those casts to fail.

The cast in question and its associated Clippy warning appear in figure 4.1. When the Solang parser is run on the `38837.sol` file from the `fuzzy-sol` repository, the code tries to cast the value 1508 to a `u8`, which is invalid.

```
warning: casting `usize` to `u8` may truncate the value
--> src/sema/expression/literals.rs:70:29
   |
70 |         ty: Type::Bytes(length as u8),
   |                               ^^^^^^^^^^^^^^^^^
   |
   = help: for further information visit
https://rust-lang.github.io/rust-clippy/master/index.html#cast_possible_truncation
```

Figure 4.1: An example warning produced by the `clippy::cast_possible_truncation` lint

```
contract TokenResolver {
    bytes4 constant INTERFACE_META_ID = 0x01ffc9a7;
    bytes4 constant ADDR_INTERFACE_ID = 0x3b3b57de;
    bytes4 constant ABI_INTERFACE_ID = 0x2203ab56;
    bytes32 constant ROOT_NODE =
0x637f12e7cd6bed65ecee34d35868279778fc56c3e5e951f46b801fb78a2d26;
    bytes TOKEN_JSON_ABI = '[{"constant":true,"inputs":[],"name":"name","output...';
```

Figure 4.2: `fuzzy-sol/Messi-Q_Smart-Contract-Dataset/38837.sol#L20-L25` (The highlighted constant is 1,508 characters long.)

Exploit Scenario

Eve discovers the invalid cast in figure 4.1. She writes a Solana program that appears to be safe, in part because of a large constant it contains. However, Eve's contract actually uses only the first 256 bytes of that constant. Eve performs a rug pull attack and uses the bug to claim plausible deniability.

Recommendations

Short term, address the bug in figure 4.1 by adding an appropriate check. Doing so will ensure that parts of large constants are not ignored.

Long term, review all casts in the codebase. For each, either add an appropriate check or document why it is guaranteed to be safe. Doing so will reduce the likelihood of bugs similar to the one in figure 4.1.

5. Rational evaluation code always returns errors

Severity: Informational

Difficulty: Low

Type: Error Reporting

Finding ID: TOB-SOLANG-5

Target: sema/expression/arithmetic.rs, sema/eval.rs

Description

The code in figure 5.1 always produces an error. Unnecessary error messages could cause user fatigue, which could make users more likely to ignore necessary errors.

Specifically, the code in figure 5.1 constructs an `Expression::Equal`. That expression is passed to the `eval_const_rational` function, which executes specific code depending on the type of expression passed to it. `eval_const_rational` does not have a case to handle `Expression::Equal`, so it returns an error.

```
let expr = Expression::Equal {
    loc: *loc,
    left: Box::new(left.cast(&left.loc(), &ty, true, ns, diagnostics)?),
    right: Box::new(right.cast(&right.loc(), &ty, true, ns, diagnostics)?),
};

if ty.is_rational() {
    if let Err(diag) = eval_const_rational(&expr, ns) {
        diagnostics.push(diag);
    }
}
```

Figure 5.1: [sema/expression/arithmetic.rs#L612-L622](#)

```
/// Resolve an expression where a compile-time constant(rational) is expected
pub fn eval_const_rational(
    expr: &Expression,
    ns: &Namespace,
) -> Result<(pt::Loc, BigRational), Diagnostic> {
    match expr {
        ...
        _ => Err(Diagnostic::error(
            expr.loc(),
            "expression not allowed in constant rational number
            expression".to_string(),
        )),
    }
}
```

Figure 5.2: [sema/eval.rs#L167-L252](#)

Recommendations

Short term, correct the code in figure 5.1 so that an error is returned only when necessary. Unnecessary error messages could cause user fatigue, which could make users more likely to ignore necessary errors.

Long term, develop additional tests for the sema module. Try to develop tests that exercise both “happy” (successful) and “sad” (failing) paths. Doing so could help to expose bugs like this one.

6. Code duplication

Severity: Informational

Difficulty: High

Type: Patching

Finding ID: TOB-SOLANG-6

Target: Various source files

Description

The Solang parser contains significant code duplication. Duplicate code can result in incomplete fixes or inconsistent behavior (e.g., if the code is modified in one location but not in all).

For example, the buggy code shown in figure 5.1 in finding [TOB-SOLANG-5](#) appears nearly verbatim in the following locations:

- [sema/expression/resolve_expression.rs#L161–L173](#)
- [sema/expression/resolve_expression.rs#L202–L214](#)
- [sema/expression/resolve_expression.rs#L242–L254](#)
- [sema/expression/resolve_expression.rs#L282–L294](#)

A second example concerns the handling of base58 constants. Code for converting such constants appears in both `sema/expression/literals.rs` (figure 6.1, left) and `sema/types.rs` (figure 6.1, right). Since this is an area where the Solang parser behaves differently than `solc`, it warrants special attention. Consolidating the code into a single function will help to ensure that correct behavior is exhibited wherever the code is used.

```
match address.from_base58() {
    Ok(v) => {
        if v.len() != ns.address_length {

diagnostics.push(Diagnostic::error(
            *loc,
            format!(
                "address literal {}
incorrect length of {}",
                address,
                v.len()
            ),
        ));
        Err(())
    }
}
```

```
match string.from_base58() {
    Ok(v) => {
        if v.len() != ns.address_length {

ns.diagnostics.push(Diagnostic::error(
            loc,
            format!(
                "address literal {}
incorrect length of {}",
                string,
                v.len()
            ),
        ));
    } else {
    }
}
```

<pre> } else { Ok(Expression::NumberLiteral { loc: *loc, ty: Type::Address(false), value: BigInt::from_bytes_be(Sign::Plus, &v), }) } } Err(FromBase58Error::InvalidBase58Length) => { diagnostics.push(Diagnostic::error(*loc, format!("address literal {address} invalid base58 length"),)); Err(()) } Err(FromBase58Error::InvalidBase58Charact er(ch, pos)) => { let mut loc = *loc; if let pt::Loc::File(_, start, end) = &mut loc { *start += pos; *end = *start; } diagnostics.push(Diagnostic::error(loc, format!("address literal {address} invalid character '{ch}'",)); Err(()) } } </pre>	<pre> seen_program_id = Some(note.loc); ns.contracts[contract_no].program_id = Some(v); } } Err(FromBase58Error::InvalidBase58Length) => { ns.diagnostics.push(Diagnostic::error(loc, format!("address literal {string} invalid base58 length"),)); } Err(FromBase58Error::InvalidBase58Charact er(ch, pos)) => { if let pt::Loc::File(_, start, end) = &mut loc { *start += pos + 1; // location includes quotes *end = *start; } ns.diagnostics.push(Diagnostic::error(loc, format!("address literal {string} invalid character '{ch}'",)); } } </pre>
---	---

Figure 6.1: *sema/expression/literals.rs#L270-L309* (left);
sema/types.rs#L433-L466 (right, with some whitespace added)

Exploit Scenario

Alice, a Solang compiler developer, is asked to fix a bug in the handling of base58 constants. Alice does not realize that the bug exists in an area other than `sema/expression/literals.rs`. Eve discovers that the bug is not fixed in `sema/types.rs` and exploits it.

Recommendations

Short term, take the following steps:

- Ensure that the bug shown in figure 5.1 is fixed in the locations listed above. Unnecessary error messages could cause user fatigue, which could make users more likely to ignore necessary errors.
- Consolidate the code for handling base58 constants into a single function. Doing so will ensure that fixes for bugs affecting such code are applied wherever the code is used.
- Review the codebase for other instances of duplicate code, and take similar code consolidation measures where possible.

Long term, adopt code practices that discourage code duplication. Doing so will help to prevent this problem from recurring.

7. Risk of arithmetic underflow in lexer

Severity: Low

Difficulty: Low

Type: Data Validation

Finding ID: TOB-SOLANG-7

Target: solang-parser/src/lexer.rs

Description

The Solang compiler's lexer contains an arithmetic operation that could underflow. This underflow can cause a panic in the compiler.

A debug build of the compiler panics on the line highlighted in figure 7.1. A release build panics on the line highlighted in figure 7.2.

```
fn parse_number(
    &mut self,
    start: usize,
    end: usize,
    ch: char,
) -> Result<(usize, Token<'input>, usize), LexicalError> {
    let mut is_rational = false;
    ...
    if ch == '.' {
        is_rational = true;
        start -= 1;
    }
    ...
    if is_rational {
        end_before_rational = start - 1;
        rational_start = start + 1;
    }
}
```

Figure 7.1: solang-parser/src/lexer.rs#L560-L622

```
let integer = &self.input[start..=end_before_rational];
```

Figure 7.2: solang-parser/src/lexer.rs#L671

```
#[test]
fn parse_number() {
    let mut comments = Vec::new();
    let mut errors = Vec::new();

    let _ = Lexer::new(".9", 0, &mut comments, &mut errors)
        .collect::<Vec<Result<(usize, Token, usize), LexicalError>>>();
}
```

```
}
```

Figure 7.3: A test that triggers the arithmetic underflow in the code in figure 7.1

Exploit Scenario

Alice tries to compile her code using the Solang compiler. The compiler crashes without producing any useful diagnostics.

Recommendations

Short term, eliminate the arithmetic underflow bug that could occur in the lexer. At a minimum, this will eliminate a panic that could occur in the compiler.

Long term, incorporate fuzzing into the CI process. Doing so could help to uncover similar bugs.

8. Excessive use of allow for lalrpop-generated code

Severity: Informational

Difficulty: High

Type: Testing

Finding ID: TOB-SOLANG-8

Target: solang-parser/src/lib.rs

Description

The Solang parser hides all Clippy warnings in the code generated from the `lalrpop` definitions. If warnings are hidden, serious problems caused by future updates to `lalrpop` or the definitions may not be detected.

```
#[allow(clippy::all)]
mod solidity {
    include!(concat!(env!("OUT_DIR"), "/solidity.rs"));
}
```

Figure 8.1: *solang-parser/src/lib.rs#L19-L22*

Recommendations

Short term, selectively allow the Clippy warnings in the generated code. This will ensure that new problems will be flagged by running Clippy.

Long term, use the `allow-all` directive sparingly, as it can hide problems that would otherwise be detected by Clippy.

9. No explicit tests for operator precedence

Severity: Informational

Difficulty: Low

Type: Testing

Finding ID: TOB-SOLANG-9

Target: The solang-parser crate

Description

The Solang parser does not have tests to verify the correct implementation of operator precedence by the `lalrpop` definitions. The expression of operator precedence using the `lalrpop` syntax is somewhat convoluted and, therefore, error-prone. While we did not uncover any problems related to this during the engagement, future updates to the definitions could cause precedence issues.

Note that, during testing, all changes that we made to operator precedence caused tests to fail. However, because those tests were actually testing properties other than those related to operator precedence, the resulting error messages did not indicate that this was the underlying problem. That is, the lack of explicit tests of operator precedence makes it hard to find the root cause of such test failures.

```
Precedence4: Expression = {
  <a:@L> <l:Precedence4> "*" <r:Precedence3> <b:@R> =>
  Expression::Multiply(Loc::File(file_no, a, b), Box::new(l), Box::new(r)),
  <a:@L> <l:Precedence4> "/" <r:Precedence3> <b:@R> =>
  Expression::Divide(Loc::File(file_no, a, b), Box::new(l), Box::new(r)),
  <a:@L> <l:Precedence4> "%" <r:Precedence3> <b:@R> =>
  Expression::Modulo(Loc::File(file_no, a, b), Box::new(l), Box::new(r)),
  Precedence3,
}

Precedence3: Expression = {
  <a:@L> <l:Precedence2> "***" <r:Precedence3> <b:@R> =>
  Expression::Power(Loc::File(file_no, a, b), Box::new(l), Box::new(r)),
  Precedence2,
}
```

Figure 9.1: An example of a change to operator precedence, from `*/%` to `**`

Recommendations

Short term, add tests verifying the correct implementation of operator precedence. This will add a safeguard to any future changes to the precedence definitions.

Long term, strive to cover as much of the parser and semantic definitions with tests as possible.

A. Vulnerability Categories

The following tables describe the vulnerability categories, severity levels, and difficulty levels used in this document.

Vulnerability Categories	
Category	Description
Access Controls	Insufficient authorization or assessment of rights
Auditing and Logging	Insufficient auditing of actions or logging of problems
Authentication	Improper identification of users
Configuration	Misconfigured servers, devices, or software components
Cryptography	A breach of system confidentiality or integrity
Data Exposure	Exposure of sensitive information
Data Validation	Improper reliance on the structure or values of data
Denial of Service	A system failure with an availability impact
Error Reporting	Insecure or insufficient reporting of error conditions
Patching	Use of an outdated software package or library
Session Management	Improper identification of authenticated users
Testing	Insufficient test methodology or test coverage
Timing	Race conditions or other order-of-operations flaws
Undefined Behavior	Undefined behavior triggered within the system

Severity Levels	
Severity	Description
Informational	The issue does not pose an immediate risk but is relevant to security best practices.
Undetermined	The extent of the risk was not determined during this engagement.
Low	The risk is small or is not one the client has indicated is important.
Medium	User information is at risk; exploitation could pose reputational, legal, or moderate financial risks.
High	The flaw could affect numerous users and have serious reputational, legal, or financial implications.

Difficulty Levels	
Difficulty	Description
Undetermined	The difficulty of exploitation was not determined during this engagement.
Low	The flaw is well known; public tools for its exploitation exist or can be scripted.
Medium	An attacker must write an exploit or will need in-depth knowledge of the system.
High	An attacker must have privileged access to the system, may need to know complex technical details, or must discover other weaknesses to exploit this issue.

B. Non-Security-Related Findings

The following recommendations are not associated with specific vulnerabilities. However, implementing them may enhance code readability and may prevent the introduction of vulnerabilities in the future.

- Adopt an API that produces diagnostics consistently. In some places, diagnostics are **pushed by the callee**:

```
if let Some(loc) = call_args_loc {
    diagnostics.push(Diagnostic::error(
        loc,
        "call arguments not permitted for internal calls".to_string(),
    ));
}
```

And in some places, diagnostics are **returned and pushed by the caller**:

```
if let Err(diag) = eval_const_rational(&expr, ns) {
    diagnostics.push(diag);
}
```

The mix of styles can be jarring for readers of the code.

- Separate **lexertest** into multiple tests:

```
#[test]
fn lexertest() {
    ... // Nearly 500 lines
}
```

Using a single, monolithic test can make it difficult to determine the cause of errors and can hamper development.