



Solang Code Generation, Part 1

Security Assessment (Summary Report)

November 17, 2023

Prepared for:

Sean Young

Solana Labs

Prepared by: **Anders Helsing**

About Trail of Bits

Founded in 2012 and headquartered in New York, Trail of Bits provides technical security assessment and advisory services to some of the world's most targeted organizations. We combine high-end security research with a real-world attacker mentality to reduce risk and fortify code. With 100+ employees around the globe, we've helped secure critical software elements that support billions of end users, including Kubernetes and the Linux kernel.

We maintain an exhaustive list of publications at <https://github.com/trailofbits/publications>, with links to papers, presentations, public audit reports, and podcast appearances.

In recent years, Trail of Bits consultants have showcased cutting-edge research through presentations at CanSecWest, HCSS, Devcon, Empire Hacking, GrrCon, LangSec, NorthSec, the O'Reilly Security Conference, PyCon, REcon, Security BSides, and SummerCon.

We specialize in software testing and code review projects, supporting client organizations in the technology, defense, and finance industries, as well as government entities. Notable clients include HashiCorp, Google, Microsoft, Western Digital, and Zoom.

Trail of Bits also operates a center of excellence with regard to blockchain security. Notable projects include audits of Algorand, Bitcoin SV, Chainlink, Compound, Ethereum 2.0, MakerDAO, Matic, Uniswap, Web3, and Zcash.

To keep up to date with our latest news and announcements, please follow [@trailofbits](#) on Twitter and explore our public repositories at <https://github.com/trailofbits>. To engage us directly, visit our "Contact" page at <https://www.trailofbits.com/contact>, or email us at info@trailofbits.com.

Trail of Bits, Inc.

228 Park Ave S #80688

New York, NY 10003

<https://www.trailofbits.com>

info@trailofbits.com

Notices and Remarks

Copyright and Distribution

© 2023 by Trail of Bits, Inc.

All rights reserved. Trail of Bits hereby asserts its right to be identified as the creator of this report in the United Kingdom.

This report is considered by Trail of Bits to be public information; it is licensed to Solana under the terms of the project statement of work and has been made public at Solana's request. Material within this report may not be reproduced or distributed in part or in whole without the express written permission of Trail of Bits.

The sole canonical source for Trail of Bits publications is the [Trail of Bits Publications page](#). Reports accessed through any source other than that page may have been modified and should not be considered authentic.

Test Coverage Disclaimer

All activities undertaken by Trail of Bits in association with this project were performed in accordance with a statement of work and agreed upon project plan.

Security assessment projects are time-boxed and often reliant on information that may be provided by a client, its affiliates, or its partners. As a result, the findings documented in this report should not be considered a comprehensive list of security issues, flaws, or defects in the target system or codebase.

Trail of Bits uses automated testing techniques to rapidly test the controls and security properties of software. These techniques augment our manual security review work, but each has its limitations: for example, a tool may not generate a random edge case that violates a property or may not fully complete its analysis during the allotted time. Their use is also limited by the time and resource constraints of a project.

Table of Contents

About Trail of Bits	1
Notices and Remarks	2
Table of Contents	3
Executive Summary	4
Project Summary	6
Project Goals	7
Project Targets	8
Project Coverage	9
Summary of Findings	10
Detailed Findings	12
1. Incorrect pointer comparison	12
2. Mixed uses of 32-bit and 64-bit values	14
3. The tests for solana.c cannot be built	16
4. Hard-coded number of accounts in SolParameters	17
5. Out-of-bounds read in the entrypoint function	18
6. Out-of-bounds read in the external_call function	20
7. Out-of-bounds write in the external_call function	22
8. Missing checks ensuring the first account passed to the contract is the state	24
10. Incorrect account deserialization	25
12. Array size changes can result in null pointer dereferences	27
13. Large array accesses cause an instruction to fail	29
14. Array access trips a llvm::checkGEPTYPE assertion	30
A. Vulnerability Categories	31
B. Non-Security-Related Findings	33

Executive Summary

Engagement Overview

Solana Labs engaged Trail of Bits to review the security of the Solang emit crate and the Solana standard library C code. From April 24 to May 5, 2023, one consultant conducted a security review of the client-provided source code, with two person-weeks of effort. Details of the project's timeline, test targets, and coverage are provided in subsequent sections of this report.

Project Scope

Our testing efforts were focused on the identification of flaws that could result in a compromise of confidentiality, integrity, or availability of the target system. We conducted this audit with full knowledge of the target system, including access to the source code and documentation. We performed static and dynamic testing of the target system and its codebase, using both automated and manual processes.

Summary of Findings

The audit uncovered issues of high severity that could impact system confidentiality, integrity, or availability. A summary of the findings is provided on the next page.

Summary of Recommendations

In addition to addressing the findings in this report, Trail of Bits recommends taking the following steps:

- Update and increase the test coverage for the C code in the `stdlib` folder, and run the tests with address sanitization and undefined behavior sanitization enabled. Incorporate the tests into the CI process to ensure they are run regularly.
- While the architecture of emitting LLVM IR code for both the Solana and Substrate target is elegant, and allows for a relatively high degree of common code between the targets, it comes with a price.

Creating loads and stores to application memory is intrinsic to the process of generating code using LLVM IR. However, for each such operation, the correct address needs to be calculated. The emit crate does this by calling the `build_gep` function from the `inkwell` crate. According to Inkwell's documentation, "GEP is very likely to segfault if indexes are used incorrectly, and is therefore an unsafe function." The code under test contains 78 calls to the function, each of which needs to correctly calculate the address in order to avoid memory corruption.

The complexity of the process of emitting LLVM IR, coupled with the specialized

traits needed for both supported targets, makes verification of the correctness of the generated code difficult, which increases the risk that bugs are not uncovered.

For the Solana target, smart contracts are typically developed in Rust, relying on the Solana Rust API or the Anchor framework. Among the benefits of using Rust is the guarantee of memory safety, and by leveraging the Anchor framework, the security considerations of contracts can be simplified.

With this in mind, our recommendation is to investigate the feasibility of drawing on these benefits by generating Rust code (to be compiled into the eBPF ELF binary) from the AST. This process could replace the current process of emitting LLVM IR altogether, or it could be done alongside it to aid in the verification of the generated code.

The following tables provide the number of findings by severity and category.

EXPOSURE ANALYSIS

<i>Severity</i>	<i>Count</i>
High	3
Medium	0
Low	0
Informational	5
Undetermined	4

CATEGORY BREAKDOWN

<i>Category</i>	<i>Count</i>
Authentication	1
Testing	1
Undefined Behavior	10

Project Summary

Contact Information

The following manager was associated with this project:

Jeff Braswell, Project Manager
jeff.braswell@trailofbits.com

The following engineering director was associated with this project:

Josselin Feist, Engineering Director, Blockchain
josselin.feist@trailofbits.com

The following engineer was associated with this project:

Anders Helsing, Consultant
anders.helsing@trailofbits.com

Project Timeline

The significant events and milestones of the project are listed below.

Date	Event
April 25, 2023	Project kickoff call
May 1, 2023	Delivery of report draft
May 8, 2023	Report readout meeting
November 17, 2023	Delivery of final report

Project Goals

The engagement was scoped to provide a security assessment of the Solang `emit` crate and Solang standard library C code. Specifically, we sought to answer the following non-exhaustive list of questions:

- Are there issues related to memory safety in the `stdlib` C code?
- Is the received contract data correctly decoded into account data, program input, and contract storage?
- Does the emitted LLVM IR accurately represent the intention of the Solidity contract code?
- Are there any security issues related to how contract storage is modeled?
- Are the existing testing strategies sufficient, or could they be extended?

Project Targets

The engagement involved a review and testing of the following targets.

Solang stdlib

Repository	https://github.com/hyperledger/solang/tree/main/stdlib
Version	3910219ed0d536dfdd017f6eb4b2828e75e99855
Type	C
Platform	Solana

Solang emit Crate

Repository	https://github.com/hyperledger/solang/tree/main/src/emit
Version	3910219ed0d536dfdd017f6eb4b2828e75e99855
Type	Rust
Platform	Solana

Solang codegen Crate

Repository	https://github.com/hyperledger/solang/tree/main/src/codegen
Version	3910219ed0d536dfdd017f6eb4b2828e75e99855
Type	Rust
Platform	Solana

Project Coverage

This section provides an overview of the analysis coverage of the review, as determined by our high-level engagement goals. Our approaches included the following:

- **Code review:**
 - We reviewed the Solana-relevant C code in the `stdlib` folder.
 - We reviewed the complete code path from the contract `entrypoint` function to the function dispatch.
 - We reviewed the `emit` crate, with a focus on code relevant to Solana.
- **Fuzzing:**
 - We built a limited fuzzer for the custom Solana storage heap, simulating extended heap operations of different sizes.

Coverage Limitations

Because of the time-boxed nature of testing work, it is common to encounter coverage limitations. Due to time constraints during this project, we could not achieve complete coverage of the `emit` crate, which may warrant further review:

1. We did not identify the root cause of three bugs in the generated LLVM IR code.
2. We performed only a cursory review of the code other than the contract storage.

Summary of Findings

The table below summarizes the findings of the review, including type and severity details.

ID	Title	Type	Severity
1	Incorrect pointer comparison	Undefined Behavior	Undetermined
2	Mixed uses of 32-bit and 64-bit values	Undefined Behavior	Undetermined
3	The tests for solana.c cannot be built	Testing	Informational
4	Hard-coded number of accounts in SolParameters	Undefined Behavior	High
5	Out-of-bounds read in the entrypoint function	Undefined Behavior	Informational
6	Out-of-bounds read in the external_call function	Undefined Behavior	High
7	Out-of-bounds write in the external_call function	Undefined Behavior	High
8	Missing checks ensuring the first account passed to the contract is the state	Authentication	Undetermined
10	Incorrect account deserialization	Undefined Behavior	Undetermined
12	Array size changes can result in null pointer dereferences	Undefined Behavior	Informational
13	Large array accesses cause an instruction to fail	Undefined Behavior	Informational

14	Array access trips a <code>llvm::checkGEPType</code> assertion	Undefined Behavior	Informational
----	--	--------------------	---------------

Detailed Findings

1. Incorrect pointer comparison

Severity: Undetermined

Difficulty: High

Type: Undefined Behavior

Finding ID: TOB-SOLEMIT-1

Target: solang/stdlib/stdlib.c

Description

The `vector_new` function receives a pointer to initial values with which to populate the vector. Our impression is that, because the zero address is valid in WASM, the function casts the pointer to an `int` and compares it to `-1` to decide whether it should be used to populate the vector. WASM is outside of the scope of this engagement, but it should be noted that dereferencing a null pointer in C is undefined behavior.

However, on Solana, the architecture is 64-bit, and casting the pointer to `int` and comparing it to `-1` has the effect that if the pointer's lower 32 bits has the value `0xffffffff`, the initialization will be omitted.

```
129 // Create a new vector. If initial is -1 then clear the data. This is done
    since a null pointer valid in wasm
130 struct vector *vector_new(uint32_t members, uint32_t size, uint8_t *initial)
131 {
132     struct vector *v;
133     size_t size_array = members * size;
134
135     v = __malloc(sizeof(*v) + size_array);
136     v->len = members;
137     v->size = members;
138
139     uint8_t *data = v->data;
140
141     if ((int)initial != -1)
```

Figure 1.1: `stdlib/stdlib.c#L129-L141`

Exploit Scenario

Mallory notices this issue and finds a contract that is affected by it. She proceeds to operate the contract outside of its intended use by triggering the bug.

Recommendations

Short term, have the function cast the value to a type with a width equal to the pointer side for the target (e.g., `ptrdiff_t`). This will ensure that the pointer is compared to the full 64-bit `-1` value on Solana, which is not a legal pointer value.

Long term, use conditional compile structures to create different implementations for the targets in order to handle the architectural differences. This will reduce the code complexity by removing the need to come up with technical workarounds that fit both targets.

2. Mixed uses of 32-bit and 64-bit values

Severity: Undetermined

Difficulty: High

Type: Undefined Behavior

Finding ID: TOB-SOLEMIT-2

Target: solang/stdlib/

Description

In several places, values that have widths of 64 bits and 32 bits are used interchangeably.

For example, the `ka_num` variable is a `uint64_t`, but it is regularly passed to functions with an `int`-type argument (figure 2.1). This is quite harmless because, in practice, a transaction cannot contain enough accounts to overflow an `int` value.

However, the `__malloc` function uses `uint32_t` for the size argument but is called in several places with a `size_t` value. The same is true for the `len` and `size` fields of the `vector` struct. Also, the `__realloc` function implicitly truncates the size of an allocated chunk from `size_t` to `uint32_t`. Should any of these values be larger than the maximum value for `uint32_t`, memory corruption will occur.

```
111     return sol_invoke_signed_c(&instruction, params->ka, params->ka_num, NULL,
0);
```

Figure 2.1: The harmless truncation of `ka_num` (`stdlib/solana.c#L111`)

```
187     struct vector *concat(uint8_t *left, uint32_t left_len, uint8_t *right,
uint32_t right_len)
188     {
189         size_t size_array = left_len + right_len;
190         struct vector *v = __malloc(sizeof(*v) + size_array);
191         v->len = size_array;
192         v->size = size_array;
```

Figure 2.2: The truncation of the `__malloc` and `vector` fields
(`stdlib/stdlib.c#L187-L192`)

```
125     void *__realloc(void *m, size_t size)
126     {
127         struct chunk *cur = m;
128
```

```

129     cur--;
130
131     struct chunk *next = cur->next;
132
133     if (next && !next->allocated && size <= (cur->length + next->length +
sizeof(struct chunk)))
134     {
135         // merge with next
136         cur->next = next->next;
137         if (cur->next)
138             cur->next->prev = cur;
139         cur->length += next->length + sizeof(struct chunk);
140         // resplit ..
141         shrink_chunk(cur, size);
142         return m;
143     }
144     else
145     {
146         // allocate new area and copy old data
147         uint32_t len = cur->length;

```

Figure 2.3: The truncation of the chunk length ([stdlib/heap.c#L125-L147](#))

Recommendations

Short term, fix all instances of inconsistent uses of type sizes that relate to memory management operations.

Long term, use type sizes consistently. Avoid implicit casts, and where a narrowing cast is intended, use an explicit cast and add a comment explaining why the cast is safe. Regularly build the `stdlib` targets with `-Wshorten-64-to-32` or `-Wconversion` to uncover new conversion issues.

3. The tests for solana.c cannot be built

Severity: Informational

Difficulty: High

Type: Testing

Finding ID: TOB-SOLEMIT-3

Target: solang/stdlib/solana.c

Description

Building the tests using `clang -DTEST -DSOL_TEST -O3 -Wall solana.c stdlib.c -o test` results in several errors. It appears that the code has been updated, but the tests have not.

Recommendations

Short term, update the test code to match the updated code.

Long term, regularly run the tests to ensure that failing tests, including build fails, are uncovered.

4. Hard-coded number of accounts in SolParameters

Severity: High

Difficulty: High

Type: Undefined Behavior

Finding ID: TOB-SOLEMIT-4

Target: solang/stdlib/solana_sdk.h

Description

The `ka` field of the `SolParameters` struct is hard-coded to be 10 elements. However, in the `sol_deserialize` function, the `ka_num` field of the same struct is set to the total number of received accounts. As a result, any function that iterates over the `ka` array using `ka_num` as the high bound (such as the functions cited in findings TOB-SOLEMIT-5, TOB-SOLEMIT-6, and TOB-SOLEMIT-7) will perform an out-of-bounds access if the number of accounts is greater than 10.

Exploit Scenario

Mallory discovers an out-of-bounds condition and finds a way to either leak back the read data or modify the contract state.

Recommendations

Short term, either have the relevant code dynamically allocate the number of entries in the `ka` array, or limit the `ka_num` field to reflect the number of entries in the array. Note that if the number of usable accounts is limited to 10, this should be reflected in the documentation.

Long term, keep the number of hard-coded arrays to a minimum.

5. Out-of-bounds read in the entrypoint function

Severity: Informational

Difficulty: High

Type: Undefined Behavior

Finding ID: TOB-SOLEMIT-5

Target: stdlib/solana.c

Description

The entrypoint function iterates over the ka array in the params struct with the ka_num field as the upper bound. This will result in an out-of-bounds read if more than 10 accounts are passed to the contract.

```
18  uint64_t
19  entrypoint(const uint8_t *input)
20  {
21      SolParameters params;
22
23      uint64_t ret = sol_deserialize(input, &params);
24      if (ret)
25      {
26          return ret;
27      }
28
29      params.ka_clock = NULL;
30      params.ka_instructions = NULL;
31
32      for (int account_no = 0; account_no < params.ka_num; account_no++)
33      {
34          const SolAccountInfo *acc = &params.ka[account_no];
35
36          if (SolPubkey_same(&clock_address, acc->key))
37          {
38              params.ka_clock = acc;
39          }
40          else if (SolPubkey_same(&instructions_address, acc->key))
41          {
42              params.ka_instructions = acc;
43          }
44      }
```

```
45
46     __init_heap();
47
48     return solang_dispatch(&params);
49 }
```

Figure 5.1: `stdlib/solana.c`#L18-L49

Because it is unlikely that the data that would be read out-of-bounds will contain the public key address of the `clock_address` or `instructions_address`, the severity of this finding is set to informational.

Recommendations

Short term, apply the short-term recommendation offered for finding **TOB-SOLEMIT-4**.

6. Out-of-bounds read in the external_call function

Severity: High

Difficulty: High

Type: Undefined Behavior

Finding ID: TOB-SOLEMIT-6

Target: stdlib/solana.c

Description

The `external_call` function iterates over the `ka` array in the `params` struct with the `ka_num` field as the upper bound. This will result in an out-of-bounds read if more than 10 accounts are passed to the contract.

```
59 // Calls an external function when 'program_id' is NULL or
60 // creates a new contract and calls its constructor.
61 uint64_t external_call(uint8_t *input, uint32_t input_len, SolPubkey
*address,
62                      SolPubkey *program_id, const SolSignerSeeds *seeds,
63                      int seeds_len, SolParameters *params)
64 {
65     SolAccountMeta metas[10];
66     SolInstruction instruction = {
67         .program_id = program_id,
68         .accounts = metas,
69         .account_len = params->ka_num,
70         .data = input,
71         .data_len = input_len,
72     };
73
74     int meta_no = 1;
75     int new_address_idx = -1;
76
77     for (int account_no = 0; account_no < params->ka_num; account_no++)
78     {
79         SolAccountInfo *acc = &params->ka[account_no];
```

Figure 6.1: `stdlib/solana.c#L59-L89`

Exploit Scenario

Mallory discovers the out-of-bounds read condition and finds a way to leak back the read data.

Recommendations

Short term, apply the short-term recommendation offered for finding **TOB-SOLEMIT-4**.

Long term, create unit tests covering the `external_call` function, and run the tests with **address sanitization** and **undefined behavior sanitization** enabled.

7. Out-of-bounds write in the external_call function

Severity: High

Difficulty: High

Type: Undefined Behavior

Finding ID: TOB-SOLEMIT-7

Target: solang/stdlib/solana.c

Description

The `external_call` function defines a `SolAccountMeta` array with a hard-coded size of 10 elements, and it iterates over this array to write accounts to it, with the `ka_num` field as the upper bound. As a result, the function will perform an out-of-bounds write if more than 10 accounts are passed to the contract.

```
59 // Calls an external function when 'program_id' is NULL or
60 // creates a new contract and calls its constructor.
61 uint64_t external_call(uint8_t *input, uint32_t input_len, SolPubkey
*address,
62                      SolPubkey *program_id, const SolSignerSeeds *seeds,
63                      int seeds_len, SolParameters *params)
64 {
65     SolAccountMeta metas[10];
66     SolInstruction instruction = {
67         .program_id = program_id,
68         .accounts = metas,
69         .account_len = params->ka_num,
70         .data = input,
71         .data_len = input_len,
72     };
73
74     int meta_no = 1;
75     int new_address_idx = -1;
76
77     for (int account_no = 0; account_no < params->ka_num; account_no++)
78     {
79         SolAccountInfo *acc = &params->ka[account_no];
80
81         // The address for the new contract should go first. Note that there
82         // may be duplicate entries, the order of those does not matter.
83         if (new_address_idx < 0 && SolPubkey_same(address, acc->key))
```

```

84     {
85         metas[0].pubkey = acc->key;
86         metas[0].is_writable = acc->is_writable;
87         metas[0].is_signer = acc->is_signer;
88         new_address_idx = account_no;
89     }
90     else
91     {
92         metas[meta_no].pubkey = acc->key;
93         metas[meta_no].is_writable = acc->is_writable;
94         metas[meta_no].is_signer = acc->is_signer;
95         meta_no += 1;
96     }
97 }

```

Figure 7.1: *stdlib/solana.c#L59-L89*

Exploit Scenario

Mallory discovers the out-of-bounds write condition and finds a way to modify the contract state.

Recommendations

Short term, apply the short-term recommendation offered for finding [TOB-SOLEMIT-4](#).

Long term, apply the long-term recommendation offered for finding [TOB-SOLEMIT-6](#).

8. Missing checks ensuring the first account passed to the contract is the state

Severity: **Undetermined**

Difficulty: **High**

Type: Authentication

Finding ID: TOB-SOLEMIT-8

Target: Various files

Description

The contract functionality implements state storage using a dedicated account, which is intended to be the first of the accounts passed to the contract. However, apart from checks to ensure that `magic` is set to the intended value, there are no checks to ensure that the first passed account is an actual state account for the contract (i.e., that it has the intended owner and address values).

Exploit Scenario

Mallory finds a contract that, if a certain state in the state account is met, uses CPI to transfer funds from a PDA vault under the contract's control. By passing in an account under her control, she fakes the necessary state and transfers funds to herself.

Recommendations

Short term, for each account used by the contract, implement checks to ensure that the address and owner of the account are the intended ones.

Long term, create integration tests, with test cases for attacks such as the one described in the exploit scenario, to ensure similar issues are uncovered if future bugs are introduced.

10. Incorrect account deserialization

Severity: **Undetermined**

Difficulty: **High**

Type: Undefined Behavior

Finding ID: TOB-SOLEMIT-10

Target: solang/stdlib/solana_sdk.h

Description

The `sol_deserialize` function iterates over received input bytes and populates the `SolParameters` array. If the number of accounts passed to the contract exceeds the hard-coded limit of the `ka` array, the input pointer is advanced, but no data is stored. However, in this case, if the `dup_info` value is not equal to `UINT8_MAX`, the pointer is not advanced to account for the 7-byte padding. As a result, the input pointer will reference incorrect data at the end of the loop, causing the `input_len`, `input`, and `program_id` variables to be incorrect.

```
241 static uint64_t sol_deserialize(  
242     const uint8_t *input,  
243     SolParameters *params)  
244 {  
245     if (NULL == input || NULL == params)  
246     {  
247         return ERROR_INVALID_ARGUMENT;  
248     }  
249     params->ka_num = *(uint64_t *)input;  
250     input += sizeof(uint64_t);  
251  
252     for (int i = 0; i < params->ka_num; i++)  
253     {  
254         uint8_t dup_info = input[0];  
255         input += sizeof(uint8_t);  
256  
257         if (i >= SOL_ARRAY_SIZE(params->ka))  
258         {  
259             if (dup_info == UINT8_MAX)  
260             {  
261                 input += sizeof(uint8_t);  
262                 input += sizeof(uint8_t);  
263                 input += sizeof(uint8_t);
```

```

264         input += 4; // padding
265         input += sizeof(SolPubkey);
266         input += sizeof(SolPubkey);
267         input += sizeof(uint64_t);
268         uint64_t data_len = *(uint64_t *)input;
269         input += sizeof(uint64_t);
270         input += data_len;
271         input += MAX_PERMITTED_DATA_INCREASE;
272         input = (uint8_t *)(((uint64_t)input + 8 - 1) & ~(8 - 1)); // padding
273         input += sizeof(uint64_t);
274     }
275     continue;
276 }
...
330 }
331
332 uint64_t data_len = *(uint64_t *)input;
333 input += sizeof(uint64_t);
334
335 params->input_len = data_len;
336 params->input = input;
337 input += data_len;
338
339 params->program_id = (SolPubkey *)input;
340 input += sizeof(SolPubkey);
341
342 return 0;
343 }

```

Figure 10.1: *stdlib/solana_sdk.h#L227-L343*

Exploit Scenario

Mallory notices the bug and finds a way to control the values of `input_len`, `input`, or `program_id` to her advantage.

Recommendations

Short term, have the function **advance the pointer to skip the padding** if `dup_info` is not equal to `UINT8_MAX`.

Long term, create unit tests for `sol_deserialize` and ensure that all code paths are covered by test cases.

12. Array size changes can result in null pointer dereferences

Severity: Informational

Difficulty: High

Type: Undefined Behavior

Finding ID: TOB-SOLEMIT-12

Target: solang/src/emit/instructions.rs

Description

Corner cases in the push operation can lead to null pointer dereferences in the emitted call to the `__realloc` function. Due to time constraints, we were unable to find the root cause of this problem.

```
165     let arr = w.vars[array].value; // Corner cases can lead to arr being null
here
166
167     let llvm_ty = bin.llvm_type(ty, ns);
168     let elem_ty = ty.array_elem();
169
170     // Calculate total size for reallocation
171     let llvm_elem_ty = bin.llvm_field_ty(&elem_ty, ns);
172     let elem_size = llvm_elem_ty
173         .size_of()
174         .unwrap()
175         .const_cast(bin.context.i32_type(), false);
176     let len = bin.vector_len(arr);
177     let new_len =
178         bin.builder
179             .build_int_add(len, bin.context.i32_type().const_int(1, false), "");
180     let vec_size = bin
181         .module
182         .get_struct_type("struct.vector")
183         .unwrap()
184         .size_of()
185         .unwrap()
186         .const_cast(bin.context.i32_type(), false);
187     let size = bin.builder.build_int_mul(elem_size, new_len, "");
188     let size = bin.builder.build_int_add(size, vec_size, "");
189
```

```

190     let realloc_size = if ns.target == Target::Solana {
191         bin.builder
192             .build_int_z_extend(size, bin.context.i64_type(), "")
193     } else {
194         size
195     };
196
197     // Reallocate and reassign the array pointer
198     let new = bin
199         .builder
200         .build_call(
201             bin.module.get_function("__realloc").unwrap(),
202             &[arr.into(), realloc_size.into()],
203             "", // Calling __realloc with null causes a null pointer dereference
204         )
205         .try_as_basic_value()
206         .left()
207         .unwrap()
208         .into_pointer_value();
209     w.vars.get_mut(array).unwrap().value = new.into();

```

Figure 12.1: solang/src/emit/instructions.rs#L165-L209

```

function foo() public pure returns (bytes memory) {
    bytes b1 = hex"42";
    bytes b2 = hex"41";

    b2.push(0x41);

    return (b1);
}

```

Figure 12.2: An example corner case that will cause a null pointer dereference

```

function foo2_func() public pure {
    uint64[] a;
    a.push(2);
}

```

Figure 12.3: Another example corner case that will cause a null pointer dereference

Recommendations

Short term, find the root problem causing the null pointer, and triage the security implications of the bug from a memory corruption perspective.

13. Large array accesses cause an instruction to fail

Severity: Informational

Difficulty: High

Type: Undefined Behavior

Finding ID: TOB-SOLEMIT-13

Target: solang/src/emit/instructions.rs

Description

Modifying the store contract's `get_values3` function to include 935 push operations on a local bytes array causes the instruction to fail, whereas 934 iterations succeed. Due to time constraints, we were unable to find the root cause of this problem.

```
function get_values3() public view returns (uint256, string memory, bytes memory,
bytes4, enum_bar) {
    bytes b2 = hex"41";

    uint i;

    for (i=0; i<935; i++)
    {
        b2.push(0x41);
    }

    return (u256, str, b2, fixedbytes, bar);
}
```

Figure 13.1: An example showing that 935 iterations cause the instruction to fail

Recommendations

Short term, find the root problem causing this issue and, if possible, put checks in place that will detect this error at compile time.

14. Array access trips a `llvm::checkGEPType` assertion

Severity: Informational

Difficulty: High

Type: Undefined Behavior

Finding ID: TOB-SOLEMIT-14

Target: `solang/src/emit/instructions.rs`

Description

A corner case function that accesses an array trips a `llvm::checkGEPType` assertion. Due to time constraints, we were unable to find the root cause of this problem.

```
function foo_func3() public view returns (uint256) {
    uint64[10] a;

    a[9] = 0x41;
    a.push(2);

    return (a[9]);
}
```

Figure 14.1: An example function that trips the `llvm::checkGEPType` assertion

```
solang:
/home/runner/work/solang-llvm/solang-llvm/llvm-project/llvm/include/llvm/IR/Instructions.h:922: llvm::Type* llvm::checkGEPType(llvm::Type*): Assertion `Ty && "Invalid GetElementPtrInst indices for type!"' failed.
Aborted (core dumped)
```

Figure 14.2: The error generated by compiling the function

Recommendations

Short term, find the root problem causing this issue and mitigate it.

A. Vulnerability Categories

The following tables describe the vulnerability categories, severity levels, and difficulty levels used in this document.

Vulnerability Categories	
Category	Description
Access Controls	Insufficient authorization or assessment of rights
Auditing and Logging	Insufficient auditing of actions or logging of problems
Authentication	Improper identification of users
Configuration	Misconfigured servers, devices, or software components
Cryptography	A breach of system confidentiality or integrity
Data Exposure	Exposure of sensitive information
Data Validation	Improper reliance on the structure or values of data
Denial of Service	A system failure with an availability impact
Error Reporting	Insecure or insufficient reporting of error conditions
Patching	Use of an outdated software package or library
Session Management	Improper identification of authenticated users
	Insufficient test methodology or test coverage
Timing	Race conditions or other order-of-operations flaws
Undefined Behavior	Undefined behavior triggered within the system

Severity Levels	
Severity	Description
Informational	The issue does not pose an immediate risk but is relevant to security best practices.
Undetermined	The extent of the risk was not determined during this engagement.
Low	The risk is small or is not one the client has indicated is important.
Medium	User information is at risk; exploitation could pose reputational, legal, or moderate financial risks.
High	The flaw could affect numerous users and have serious reputational, legal, or financial implications.

Difficulty Levels	
Difficulty	Description
Undetermined	The difficulty of exploitation was not determined during this engagement.
Low	The flaw is well known; public tools for its exploitation exist or can be scripted.
Medium	An attacker must write an exploit or will need in-depth knowledge of the system.
High	An attacker must have privileged access to the system, may need to know complex technical details, or must discover other weaknesses to exploit this issue.

B. Non-Security-Related Findings

The following finding is not associated with specific vulnerabilities. However, addressing it will enhance code readability and may prevent the introduction of vulnerabilities in the future.

- The **documentation** of the data field of the AccountInfo struct states that “This field can be modified, but use with caution”; however, it is marked as read-only in **the code**.