



# Solang Code Generation

Security Assessment (Summary Report)

November 17, 2023

*Prepared for:*

**Sean Young**

Solana Labs

*Prepared by:* **Samuel Moelius and Vara Prasad Bandaru**

# About Trail of Bits

---

Founded in 2012 and headquartered in New York, Trail of Bits provides technical security assessment and advisory services to some of the world's most targeted organizations. We combine high-end security research with a real-world attacker mentality to reduce risk and fortify code. With 100+ employees around the globe, we've helped secure critical software elements that support billions of end users, including Kubernetes and the Linux kernel.

We maintain an exhaustive list of publications at <https://github.com/trailofbits/publications>, with links to papers, presentations, public audit reports, and podcast appearances.

In recent years, Trail of Bits consultants have showcased cutting-edge research through presentations at CanSecWest, HCSS, Devcon, Empire Hacking, GrrCon, LangSec, NorthSec, the O'Reilly Security Conference, PyCon, REcon, Security BSides, and SummerCon.

We specialize in software testing and code review projects, supporting client organizations in the technology, defense, and finance industries, as well as government entities. Notable clients include HashiCorp, Google, Microsoft, Western Digital, and Zoom.

Trail of Bits also operates a center of excellence with regard to blockchain security. Notable projects include audits of Algorand, Bitcoin SV, Chainlink, Compound, Ethereum 2.0, MakerDAO, Matic, Uniswap, Web3, and Zcash.

To keep up to date with our latest news and announcements, please follow [@trailofbits](#) on Twitter and explore our public repositories at <https://github.com/trailofbits>. To engage us directly, visit our "Contact" page at <https://www.trailofbits.com/contact>, or email us at [info@trailofbits.com](mailto:info@trailofbits.com).

## Trail of Bits, Inc.

228 Park Ave S #80688

New York, NY 10003

<https://www.trailofbits.com>

[info@trailofbits.com](mailto:info@trailofbits.com)

# Notices and Remarks

---

## Copyright and Distribution

© 2023 by Trail of Bits, Inc.

All rights reserved. Trail of Bits hereby asserts its right to be identified as the creator of this report in the United Kingdom.

This report is considered by Trail of Bits to be public information; it is licensed to Solana Labs under the terms of the project statement of work and has been made public at Solana Labs' request. Material within this report may not be reproduced or distributed in part or in whole without the express written permission of Trail of Bits.

The sole canonical source for Trail of Bits publications is the [Trail of Bits Publications page](#). Reports accessed through any source other than that page may have been modified and should not be considered authentic.

## Test Coverage Disclaimer

All activities undertaken by Trail of Bits in association with this project were performed in accordance with a statement of work and agreed upon project plan.

Security assessment projects are time-boxed and often reliant on information that may be provided by a client, its affiliates, or its partners. As a result, the findings documented in this report should not be considered a comprehensive list of security issues, flaws, or defects in the target system or codebase.

Trail of Bits uses automated testing techniques to rapidly test the controls and security properties of software. These techniques augment our manual security review work, but each has its limitations: for example, a tool may not generate a random edge case that violates a property or may not fully complete its analysis during the allotted time. Their use is also limited by the time and resource constraints of a project.

# Table of Contents

---

<b>About Trail of Bits</b>	<b>1</b>
<b>Notices and Remarks</b>	<b>2</b>
<b>Table of Contents</b>	<b>3</b>
<b>Executive Summary</b>	<b>4</b>
<b>Project Summary</b>	<b>5</b>
<b>Project Goals</b>	<b>6</b>
<b>Project Targets</b>	<b>7</b>
<b>Summary of Findings</b>	<b>8</b>
<b>Detailed Findings</b>	<b>10</b>
1. Use of dependency with open RUSTSEC advisory	10
2. Use of outdated dependencies	12
3. Insufficient test coverage	14
4. Tests do not pass with latest stable Rust version	16
5. Strength reduction does not properly handle undefined variables	17
6. Solang fails to compile certain valid recursive structures	20
7. Monolithic test	22
8. Optimizations hide unused undefined variables in contracts	24
9. Solang-compiled contracts can have multiple storage accounts	25
10. An attacker can reinitialize a Solang contract	27
11. Compiler does not verify the developer-specified size for the data account	29
12. The bump is not guaranteed to be at the end of seeds array	31
13. Appending state variables to Solang contracts affects their storage layout	33
14. Fallback function does not verify the data account's magic value	34
<b>A. Vulnerability Categories</b>	<b>37</b>
<b>B. Non-Security-Related Findings</b>	<b>39</b>

# Executive Summary

---

## Engagement Overview

Solana Labs engaged Trail of Bits to review the security of Solang's codegen module, specifically of how it generates Solana code.

A team of two consultants conducted the review from June 23 to July 12, 2023, for a total of four engineer-weeks of effort. With full access to the source code and documentation, we performed static and dynamic testing of the codebase, using automated and manual processes.

## Observations and Impact

As discussed in finding [TOB-SOLCG-3](#), there are no tests to verify that unoptimized and optimized code behave the same. During the project kickoff call, the Solang team described an improperly applied optimization as a "worst case scenario." Having tests to help verify the optimization passes' correctness is the best way to defend against such possibilities. Hence, we highly recommend that such tests be added.

The following tables provide the number of findings by severity and category.

### EXPOSURE ANALYSIS

<i>Severity</i>	<i>Count</i>
High	4
Medium	1
Low	3
Informational	6
Undetermined	0

### CATEGORY BREAKDOWN

<i>Category</i>	<i>Count</i>
Data Validation	6
Patching	2
Testing	3
Undefined Behavior	3

# Project Summary

---

## Contact Information

The following manager was associated with this project:

**Jeff Braswell**, Project Manager  
[jeff.braswell@trailofbits.com](mailto:jeff.braswell@trailofbits.com)

The following engineering director was associated with this project:

**Josselin Feist**, Engineering Director, Blockchain  
[josselin.feist@trailofbits.com](mailto:josselin.feist@trailofbits.com)

The following engineers were associated with this project:

**Samuel Moelius**, Consultant      **Vara Prasad Bandaru**, Consultant  
[samuel.moelius@trailofbits.com](mailto:samuel.moelius@trailofbits.com)      [vara.bandaru@trailofbits.com](mailto:vara.bandaru@trailofbits.com)

## Project Timeline

The significant events and milestones of the project are listed below.

Date	Event
June 15, 2023	Technical onboarding call
June 23, 2023	Pre-project kickoff call
June 30, 2023	Status update meeting #1
July 12, 2023	Delivery of report draft
July 12, 2023	Report readout meeting
July 24, 2023	Addition of finding <b>TOB-SOLCG-14</b> ; updates to <b>appendix B</b>
November 17, 2023	Delivery of final report

# Project Goals

---

The engagement was scoped to provide a security assessment of Solang's codegen module, specifically of how it generates Solana code. We sought to answer the following non-exhaustive list of questions:

- Does code emitted by the codegen module preserve the semantics of the original source code?
- Are optimizations applied under appropriate circumstances?
- Do optimizations preserve the semantics of the unoptimized code?
- Does Solang's code generation strategy introduce behavior that would be surprising to Solidity or Solana developers?

# Project Targets

---

The engagement involved a review and testing of the following target.

## Solang codegen Module

Repository	<a href="https://github.com/hyperledger/solang/tree/main/src/codegen">https://github.com/hyperledger/solang/tree/main/src/codegen</a>
Version	a84b0ad3b67a17b524ef6b7437fd4c5376833807
Types	Rust/Solidity
Platform	Solana



## Summary of Findings

The table below summarizes the findings of the review, including type and severity details.

ID	Title	Type	Severity
1	Use of dependency with open RUSTSEC advisory	Patching	Informational
2	Use of outdated dependencies	Patching	Informational
3	Insufficient test coverage	Testing	Informational
4	Tests do not pass with latest stable Rust version	Testing	Informational
5	Strength reduction does not properly handle undefined variables	Data Validation	Low
6	Solang fails to compile certain valid recursive structures	Data Validation	Low
7	Monolithic test	Testing	Informational
8	Optimizations hide unused undefined variables in contracts	Undefined Behavior	Informational
9	Solang-compiled contracts can have multiple storage accounts	Data Validation	High
10	An attacker can reinitialize a Solang contract	Data Validation	High
11	Compiler does not verify the developer-specified size for the data account	Data Validation	Medium
12	The bump is not guaranteed to be at the end of seeds array	Undefined Behavior	Low

13	Appending state variables to Solang contracts affects their storage layout	Undefined Behavior	High
14	Fallback function does not verify the data account's magic value	Data Validation	High

## Detailed Findings

### 1. Use of dependency with open RUSTSEC advisory

Severity: Informational

Difficulty: Undetermined

Type: Patching

Finding ID: TOB-SOLCG-1

Target: Cargo.lock

### Description

The borsh dependency (which the codegen module relies on) has an outstanding **RUSTSEC advisory**. A fix has been merged but apparently not released. Solang should use an updated version of borsh as soon as one is released with the fix.

The following is an excerpt from the RUSTSEC advisory:

*Affected versions of borsh cause undefined behavior when zero-sized-types (ZST) are parsed and the Copy/Clone traits are not implemented/derived. For instance if 1000 instances of a ZST are deserialized, and the ZST is not copy (this can be achieved through a singleton), then accessing/writing to deserialized data will cause a segmentation fault.*

*There is currently no way for borsh to read data without also providing a Rust type. Therefore, if not [sic] ZST are used for serialization, then you are not affected by this issue.*

A fix was merged on June 7, 2023. However, as of this writing, the fix does not appear in any release.

Note that **cargo-audit** warns about dependencies besides borsh. However, none of those dependencies are used by the codegen module.

### Exploit Scenario

Alice, a Solang developer, writes a test that uses zero sum types. Eve learns of this fact and exploits the bug on Alice's machine.

### Recommendations

Short term, watch the borsh repository and switch to a new version of borsh as soon as one is released with the fix. Doing so will help ensure that Solang developers and users do not use vulnerable dependencies.

Long term, regularly run `cargo-audit` over the codebase. Doing so will help to identify vulnerable or unmaintained dependencies.

## References

- [RUSTSEC-2023-0033: Parsing borsh messages with ZST which are not-copy/clone is unsound](#)
- [borsh-rs issue #13: BorshDeserialize can cause UB by copying zero sized objects with no safe Copy impl](#)
- [borsh-rs pull request #145: Forbid Zero-sized types from deserialization](#)

## 2. Use of outdated dependencies

Severity: **Informational**

Difficulty: **High**

Type: Patching

Finding ID: TOB-SOLCG-2

Target: Cargo.toml

### Description

Updated versions of many of the codegen module's dependencies are available. Because silent bug fixes are common, all dependencies should be periodically reviewed and updated wherever possible.

Note that some of these outdated dependencies have updated versions that are considered incompatible by Cargo; because of this, simply running cargo update will not cause them to be updated in the project's Cargo.lock file. Dependencies for which incompatible upgrades are available appear in table 2.1.

Dependency	Version currently in use	Latest version available
itertools	0.10.5 (September 18, 2022)	0.11.0 (June 22, 2023)
indexmap	1.9.3 (March 24, 2023)	2.0.0 (June 23, 2023)
anchor-syn	0.27.0 (March 8, 2023)	0.28 (June 9, 2023)

*Table 2.1: Dependencies for which incompatible upgrades are available*

Note that dependencies besides those in table 2.1 can be upgraded. However, none of those dependencies are used by the codegen module.

### Exploit Scenario

Eve learns of a vulnerability in an outdated version of a codegen dependency. Knowing that the codegen module still relies on this outdated version, Eve exploits the vulnerability.

### Recommendations

Short term, update the dependencies to their latest versions wherever possible. Verify that all unit tests pass following such updates. Document any reasons for not updating a dependency. Using out-of-date dependencies could mean critical bug fixes are missed.

Long term, regularly run `cargo upgrade --incompatible`. This will help ensure that the project stays up to date with its dependencies.

### 3. Insufficient test coverage

Severity: Informational

Difficulty: High

Type: Testing

Finding ID: TOB-SOLCG-3

Target: The tests subdirectory

#### Description

The codegen module would benefit from additional tests.

The tests most applicable to generating Solana code are the codegen and solana tests. The former tests that certain patterns are present in generated code. The latter tests that generated code behaves correctly.

Tables 3.1 and 3.2 summarize the code covered by the aforementioned tests. Note that code not covered by the codegen test may be covered by the solana test, and vice versa. Put another way, the red sections in tables 3.1 and 3.2 should not be interpreted as “not covered by any test.”

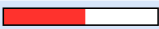
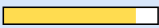


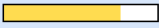
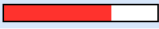
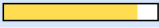
<a href="#">src/codegen</a>		53.5 %	5827 / 10900	60.8 %	192 / 316
<a href="#">src/codegen/dispatch</a>		85.7 %	766 / 894	95.0 %	19 / 20
<a href="#">src/codegen/encoding</a>		73.9 %	1936 / 2618	61.3 %	57 / 93
<a href="#">src/codegen/events</a>		0.0 %	0 / 224	0.0 %	0 / 8
<a href="#">src/codegen/solana_accounts</a>		76.3 %	425 / 557	78.6 %	11 / 14
<a href="#">src/codegen/strength_reduce</a>		70.2 %	957 / 1363	84.7 %	61 / 72
<a href="#">src/codegen/subexpression_elimination</a>		87.1 %	1451 / 1666	84.4 %	81 / 96

Table 3.1: This table shows the code covered by the codegen test. The four rightmost columns represent the percentage of lines covered, number of lines covered, percentage of functions covered, and number of functions covered.


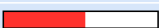

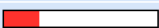
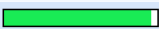
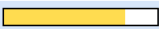
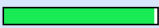
<a href="#">src/codegen</a>		72.6 %	7908 / 10900	69.6 %	220 / 316
<a href="#">src/codegen/dispatch</a>		52.6 %	470 / 894	40.0 %	8 / 20
<a href="#">src/codegen/encoding</a>		77.6 %	2031 / 2618	61.3 %	57 / 93
<a href="#">src/codegen/events</a>		23.2 %	52 / 224	37.5 %	3 / 8
<a href="#">src/codegen/solana_accounts</a>		96.1 %	535 / 557	100.0 %	14 / 14
<a href="#">src/codegen/strength_reduce</a>		78.7 %	1072 / 1363	84.7 %	61 / 72
<a href="#">src/codegen/subexpression_elimination</a>		97.9 %	1631 / 1666	86.5 %	83 / 96

Table 3.2: This table shows the code covered by the solana test. The four rightmost columns represent the percentage of lines covered, number of lines covered, percentage of functions covered, and number of functions covered.

Notably missing from the above is a test that does the following:

- Compiles a Solidity program with optimizations disabled
- Runs the resulting binary on one or more test vectors
- Compiles the same program with optimizations enabled
- Runs the resulting binary on the same set of test vectors
- Verifies that the two binaries' outputs are equal

Ideally, this test would operate on a large number of Solidity programs and would have many test vectors for each.

### **Exploit Scenario**

A bug is found in an optimization pass. The bug could have been exposed by more thorough unit or integration tests.

### **Recommendations**

Short term, add tests to compile code with and without optimizations, and verify that the resulting binaries behave similarly. Doing so will help increase confidence in the code that performs optimizations.

Long term, regularly compute and review test coverage using a tool such as [cargo-llvm-cov](#). Doing so will help ensure that the tests are relevant and that all important conditions are tested.



#### 4. Tests do not pass with latest stable Rust version

Severity: Informational

Difficulty: Undetermined

Type: Testing

Finding ID: TOB-SOLCG-4

Target: The tests subdirectory

#### Description

The tests do not pass when built with the latest version of the Rust compiler (1.70.0). To ensure that the code can benefit from compiler bug fixes, the code should be kept up to date with the latest stable Rust version.

The error message produced by running the solana test compiled with Rust 1.70.0 appears in figure 4.1.

```
thread 'solana_tests::abi_decode::decode_address' panicked at 'misaligned pointer
dereference: address must be a multiple of 0x8 but is 0x7f724841a82c',
.../solana_rbpf-0.2.38/src/interpreter.rs:270:26
```

*Figure 4.1: The error message produced by running the solana test compiled with Rust 1.70.0*

#### Exploit Scenario

Rust version 1.70.1 fixes a critical bug in the compiler. Because Solang cannot be compiled with Rust 1.70.0, Solang does not benefit from the bug fix. Eve notices this and exploits the Solang instance running on Alice's machine.

#### Recommendations

Short term, diagnose and fix all tests that do not pass when compiled with Rust 1.70.0. Doing so will allow the code to benefit from fixes to the current stable version of Rust and will ease the transition to the next version.

Long term, regularly test the code with the latest stable Rust version. Doing so will help the code to benefit from compiler bug fixes.

## 5. Strength reduction does not properly handle undefined variables

Severity: Low

Difficulty: Low

Type: Data Validation

Finding ID: TOB-SOLCG-5

Target: codegen/cfg.rs

### Description

The strength reduction optimization runs even when undefined variables are present. This can result in an assertion violation and a panic.

For example, a panic can be triggered by changing the equal sign [=] highlighted in figure 5.1 to a semicolon [;], which would make the `info` variable undefined. The panic occurs in the code in figure 5.2. Two other parts of the call chain appear in figures 5.3 and 5.4. (Several call frames that would appear between figures 5.3 and 5.4 are omitted.) Note the comments in figure 5.4, which appear not to accurately reflect the current code.

```
contract MyTest {
    // BEGIN-CHECK: MyTest::MyTest::function::test_this__uint32_address
    function test_this(uint32 i, address addr) public view returns (uint32) {
        AccountInfo info = tx.accounts[i];
        if (info.key == addr) {
            // CHECK: branchcond ((load (load (struct %info field 0))) == (arg #1)),
            block3, block4
            return 0;
        } else if (info.lamports == 90) {
```

*Figure 5.1: Changing the highlighted equal sign [=] to a semicolon [;] makes `info` undefined and causes a panic.*

*(tests/codegen\_testcases/solidity/load\_account\_info\_members.sol#L5-L12)*

```
impl Type {
    /// Default value for a type, e.g. an empty string. Some types cannot have a
    default value,
    /// for example a reference to a variable in storage.
    pub fn default(&self, ns: &Namespace) -> Option<Expression> {
        match self {
            ...
            Type::Ref(ty) => {
                assert!(matches!(ty.as_ref(), Type::Address(_)));
            }
            ...
        }
    }
}
```

*Figure 5.2: codegen/statements.rs#L1440-L1488*

```

impl Type {
    /// Default value for a type, e.g. an empty string. Some types cannot have a
    pub(super) fn expression_values(
        expr: &Expression,
        vars: &Variables,
        ns: &Namespace,
    ) -> HashSet<Value> {
        match expr {
            ...
            Expression::Undefined { ty } => {
                // If the variable is undefined, we can return the default value to
                optimize operations
                if let Some(default_expr) = ty.default(ns) {
                    return expression_values(&default_expr, vars, ns);
                }

                HashSet::new()
            }
        }
    }
}

```

Figure 5.3: *codegen/strength\_reduce/expression\_values.rs#L13–L84*

```

/// Detect undefined variables and run codegen optimizer passess
pub fn optimize_and_check_cfg(
    cfg: &mut ControlFlowGraph,
    ns: &mut Namespace,
    func_no: ASTFunction,
    opt: &Options,
) {
    reaching_definitions::find(cfg);
    if func_no != ASTFunction::None {
        // If there are undefined variables, we raise an error and don't run
        optimizations
        if undefined_variable::find_undefined_variables(cfg, ns, func_no) {
            return;
        }
    }
    if opt.constant_folding {
        constant_folding::constant_folding(cfg, ns);
    }
    if opt.vector_to_slice {
        vector_to_slice::vector_to_slice(cfg, ns);
    }
    if opt.strength_reduce {
        strength_reduce::strength_reduce(cfg, ns);
    }
}

```

Figure 5.4: *codegen/cfg.rs#L1539–L1561*

## Exploit Scenario

Alice tries to compile her code using the Solang compiler. The compiler crashes without producing any useful diagnostics.

## Recommendations

Short term, eliminate the assertion failure that can occur in the code in figure 5.2. Doing so will eliminate a panic that could occur in the codegen module.

Long term, incorporate fuzzing into the CI process. Doing so could help to reveal similar bugs.

## 6. Solang fails to compile certain valid recursive structures

Severity: Low

Difficulty: Low

Type: Data Validation

Finding ID: TOB-SOLCG-6

Target: The sema module

### Description

Solang considers structs containing multidimensional dynamic-size arrays of the structs' own type with fixed-size innermost arrays to be of infinite size; therefore, Solang fails to compile them.

Structs containing dynamic-size arrays of the structs' own type, irrespective of dimensions, should be considered to be of finite size, and compilation should be possible.

Figure 6.1 contains an example struct definition that contains a dynamic-size array of the struct's own type with a *dynamic-size* innermost array. Solang correctly considers this struct to be of finite size and successfully compiles it.

```
struct A {  
    A[][1][2] b;  
}
```

Figure 6.1: An example struct containing a dynamic-size array of the struct's own type with a dynamic-size innermost array

Figure 6.2 contains an example struct definition that contains a dynamic-size array of the struct's own type but with a *fixed-size* innermost array. Solang fails to compile it with the error "struct 'A' has infinite size."

```
struct A {  
    A[2][1][] b;  
}
```

Figure 6.2: An example struct containing a dynamic-size array of the struct's type with a fixed-sized innermost array

### Exploit Scenario

A contract contains a struct definition containing a dynamic-size array of the struct's own type with a fixed-size innermost array, similar to the definition in figure 6.2. The compiler fails with the error "struct has infinite size."

## Recommendations

Short term, correct Solang's handling of recursive structures such as the code in figure 6.2. As this code is valid Solidity, it should be accepted.

Long term, improve the tests for the compilation of recursive structs. Doing so will help to identify problems like the one described in this finding.

## 7. Monolithic test

Severity: Informational

Difficulty: High

Type: Testing

Finding ID: TOB-SOLCG-7

Target: codegen/strength\_reduce/tests.rs

### Description

The `expresson_known_bits` test is approximately 1,200 lines (figure 7.1). Large tests can prevent errors from being caught and can hamper future development.

```
#[test]
fn expresson_known_bits() {
    use crate::Target;
    use solang_parser::pt::Loc;
    ...
    ... // just under 1200 lines
    ...
    assert!(v.known_bits[0]);
    assert!(v.value[0]);
}
```

Figure 7.1: `codegen/strength_reduce/tests.rs#L29-L1230`

There are good reasons to break a large test up into multiple smaller tests.

First, if a large test fails, it could be difficult for a developer to determine the cause. More specifically, if the test fails on the  $n$ th statement, it could be difficult for the developer to determine which of the  $n-1$  preceding statements contributed to the failure.

Second, an oft overlooked benefit of tests is that they serve as documentation. However, a monolithic test detracts from this benefit. Suppose a developer wants to know how to use statement X, which happens to be on line  $n$  of the test. If  $n$  is large, it could be difficult for the developer to determine which of the preceding  $n-1$  statements were necessary to use X.

### Exploit Scenario

Alice, a Solang developer, makes a change to the code that causes the `expresson_known_bits` test to fail. The amount of time that Alice spends trying to determine the cause of the failure is more than it would have been had a smaller test failed.

## Recommendations

Short term, break the `expresson_known_bits` test up into smaller tests. This will make determining the cause of failures easier and will help streamline future development.

Long term, consider enabling Clippy's `too-many-lines` lint and setting its lint level to deny. Doing so will help limit the size of future tests.



## 8. Optimizations hide unused undefined variables in contracts

Severity: Informational

Difficulty: Low

Type: Undefined Behavior

Finding ID: TOB-SOLCG-8

Target: The codegen module

### Description

The compiler does not raise an error for contracts containing undefined variables when optimizations are enabled. As a result, developers might not be aware of incorrectness in their contracts.

The compiler runs the “remove unused variables” optimization before undefined variable detection. As a result, the “remove unused variables” optimization will remove any undefined variables that are not used, and undefined variable detection will not raise the appropriate error. However, when optimizations are disabled, undefined variables will not be removed, and the compiler will raise the undefined variable error.

For example, the share variable in the contract in figure 8.1 is undefined. The compiler would raise the undefined variable error if the contract were compiled without optimizations. However, with optimizations enabled, the contract would be compiled without any warnings or errors.

```
contract Test {  
    struct A {  
        uint256 b;  
    }  
  
    function test() public {  
        A storage share;  
        share.b = uint(10);  
    }  
}
```

*Figure 8.1: An example contract containing an undefined variable*

### Recommendations

Short term, update the implementation to run undefined variable detection before performing any optimizations.

Long term, write tests to verify the equivalence of the code compiled with and without optimizations.

## 9. Solang-compiled contracts can have multiple storage accounts

Severity: **High**

Difficulty: **Medium**

Type: Data Validation

Finding ID: TOB-SOLCG-9

Target: codegen/solana\_deploy.rs

### Description

The compiler-generated constructor code does not ensure the uniqueness of the contract's data account, which might lead to account confusion issues in which a data account different from the intended data account can be used.

The contract storage is represented using a data account. All of the state variables are stored in that account. The constructor initializes the data account by writing a magic value in the first eight bytes of the account data. This magic value is used by the contract functions to verify that the correct data account is passed, ensuring that the correct account is used for storage.

The constructor does not prevent users from creating multiple data accounts. Any user can call the constructor with a new account, and the constructor will write the same magic value to that account. The new account can be used as the storage for the contract. This allows users to create a single deployment for multiple instances of a contract, each with its own storage. All instances will have the same program ID but different data accounts.

The disadvantage of this is that the users and protocols interacting with the contract will have to ensure that the intended data account is being used by the contract (i.e., that they are interacting with the intended instance of the contract).

This approach becomes an issue when part of the contract's state is independent of the storage. For example, if a contract uses a program-derived address (PDA) account to interact with external contracts, that PDA account can be considered part of the contract's state. The PDA depends on the program ID and a list of seeds. If the seeds are static and fixed at compile time, the derived PDA will be independent of the contract's storage.

When the PDA is independent of the contract storage and depends only on the code, all instances of the contract with different storage accounts will use the same PDA account. This creates an overlap between states of different instances of the contract. An attacker can exploit this by creating a new data account with storage favorable to them and using the PDA of existing instances to perform operations and profit from them.

## Exploit Scenario

Consider a contract with the following description:

- The constructor sets the owner state variable to the caller-given account.
- The contract owns tokens using the PDA derived from the seed [ "token owner" ].
- The contract contains the `withdraw` function, which, when called by the owner and the owner is a signer, transfers tokens owned by the PDA to the owner account.

Bob, the developer, deploys the contract and calls the contract with data account A. The owner value in account A is owned by Bob. After some time, with the normal use of the contract, the PDA derived from the [ "token owner" ] seed owns \$1 million worth of tokens.

Eve, an attacker, calls the constructor with data account B. The owner value in account B is owned by Eve. Eve calls the `withdraw` function using data account B and Bob's PDA. Because the PDA does not depend on the storage, it is the same for Eve's instance as well. The `withdraw` function succeeds, and Eve steals the tokens owned by Bob.

## Recommendations

Short term, update the compiler to ensure the uniqueness of the data account for a given program ID and the contract. This can be achieved by ensuring that the data account is a PDA derived using static seeds. Alternatively, if the feature is needed, add warnings to the developer documentation explaining the risks of the current approach. Also add documentation advising external protocols and users interacting with Solang contracts to verify the data account's address.

Long term, document the design choices along with the assumptions made and perform a review to ensure that the selected design choices do not break the system invariants.

## 10. An attacker can reinitialize a Solang contract

Severity: High

Difficulty: Low

Type: Data Validation

Finding ID: TOB-SOLCG-10

Target: codegen/solana\_deploy.rs

### Description

The compiler-generated constructor code does not check whether a data account is already initialized. As a result, an attacker can call the constructor using an initialized data account and update important state variables.

The contract storage is represented using a data account. All of the state variables are stored in that account. The constructor initializes the data account by writing a magic value in the first eight bytes of the account data.

Before running the initialization routines, the constructor does not check the account's magic value and proceeds with initialization. As a result, the state variables initialized in the constructor will be updated with the initial values and the caller-provided arguments.

### Exploit Scenario

Bob, a developer, deploys the Test contract (figure 10.1). He calls the constructor and sets the owner variable to his address. After some time, with continuous use of the contract, the contract owns assets worth \$10 million.

```
contract Test {
    address owner;

    constructor(address admin) {
        owner = admin;
    }
    [...]

    function withdraw() public {
        // verify owner is signer and transfer all assets.
    }
}
```

Figure 10.1: An example contract that is vulnerable to this issue

Eve, an attacker, calls the constructor with her address as the admin variable. The constructor updates the owner variable. Eve calls the withdraw function and steals \$10 million worth of assets.

## Recommendations

Short term, update the `solana_deploy` function to add initialization checks in the constructor code.

Long term, write a reference implementation in a high-level language for every instance of compiler-generated code written using low-level codegen instructions. Review the high-level reference implementation and ensure that the low-level implementation is equivalent to the reference implementation.

## 11. Compiler does not verify the developer-specified size for the data account

Severity: **Medium**

Difficulty: **Medium**

Type: Data Validation

Finding ID: TOB-SOLCG-11

Target: `codegen/solana_deploy.rs`

### Description

The constructor does not verify the minimum size requirement for the data account while creating the account using a developer-provided value. As a result, the data account could become unusable during the use of the contract.

The constructor creates the data account if it is not given by the caller. The data account is required to have a certain minimum size. The developer can specify the data account size using the `space` annotation. The `space` value could be static, known during the compilation, or it could be dynamic, given as an argument. However, the compiler neither performs compile time checks nor adds runtime checks for the `space` value. If the developer incorrectly calculates the required size or mistakenly provides the wrong value, the created data account could have less space than required.

The minimum size is referred to as the `contract.fixed_layout_size`. It represents the size required to store the contract's fixed-size storage variables. If the data account is a size less than `fixed_layout_size`, then only the first few variables can be read or written. All operations that require reading or writing the fixed-size variables stored at the end will fail with an out-of-bounds error.

Because only some of the operations might fail, the issue may not be caught during early use of the contract, and the contract could become unusable in an intermediate state.

### Exploit Scenario

The `fixed_layout_size` value for the contract in figure 11.1 is 2048 bytes. Bob, the developer, mistakenly specifies 2000 bytes in the `space` annotation. The compiler compiles the code without raising any errors. Bob deploys the contract and calls the constructor. The constructor creates the data account with a size of 2000 bytes and initializes the account.

The deposit operations and other operations succeed without any errors. After some time, the contract accumulates assets. Bob tries to withdraw the assets using the `withdraw` function. The `withdraw` function writes to the `withdrawn` variable, which is stored after the offset 2000 in the data account. The operation fails with an out-of-bounds error, and the funds are stuck in the contract.

```

contract Test {
    address owner;
    [...]
    bool withdrawn;

    @payer(...)
    @space(2000)
    constructor(address admin) {
        owner = admin;
    }

    function deposit() public { [...] }
    function withdraw() public {
        // verify owner is signer and transfer all assets.
        // The function writes to the `withdrawn` variable.
    }
}

```

*Figure 11.1: An example contract that is vulnerable to this issue*

## Recommendations

Short term, update the `solana_deploy` function to check the `space` value during compilation if it is static and to add runtime checks to the constructor code if the `space` value is a runtime constant.

Long term, implement the compiler to be strict and perform as many checks as possible. Develop the compiler with the assumption that the developer will make mistakes and write incorrect code.

## 12. The bump is not guaranteed to be at the end of seeds array

Severity: Low

Difficulty: Medium

Type: Undefined Behavior

Finding ID: TOB-SOLCG-12

Target: codegen/solana\_deploy.rs

### Description

While constructing the seeds array using the constructor annotations, the compiler does not ensure that the bump value is placed at the end of the array. As a result, the computed account might not be a valid PDA and the contract initialization might fail.

The developer can specify the seeds and bump value for a PDA using the constructor annotations. The compiler uses the seeds in the specified order for signing the PDA account. It considers the bump value as just another seed value and includes it in the specified position.

```
for note in &func.annotations {
    match note {
        ConstructorAnnotation::Seed(seed) => {
            seeds.push(expression(seed, cfg, contract_no, None, ns, vartab, opt));
        }
        ConstructorAnnotation::Bump(bump) => {
            let expr = ast::Expression::Cast {
                loc: Loc::Codegen,
                to: Type::Slice(Type::Bytes(1).into()),
                expr: ast::Expression::BytesCast {
                    loc: Loc::Codegen,
                    to: Type::DynamicBytes,
                    from: Type::Bytes(1),
                    expr: bump.clone().into(),
                }
            }.into(),
        };
        seeds.push(expression(&expr, cfg, contract_no, None, ns, vartab, opt));
    }
    - => (),
}
```

Figure 12.1: *codegen/solana\_deploy.rs#L463-L484*

However, the bump value is expected to be the last seed, so it should be placed at the end of the array. The developer might work with the assumption that the compiler will place the



bump value at the end irrespective of the position of its annotation. If the developer places a seed annotation after the bump annotation, the order of seeds used by the compiler will be different from the order expected by the developer.

Because the order of the seeds decides the derived PDA account, the derived address will be different than expected and might not be a valid account if the bump value is not at the end of the array. The derived PDA account is needed for contract initialization. As a result, the contract might need to be redeployed after the position of the bump annotation is updated.

### Exploit Scenario

Bob, a developer, expects the seeds array for the PDA of the contract in figure 12.2 to be [ "A", "B", "x" ]. However, the seeds used by the compiler for the PDA will be [ "A", "x", "B" ]. Bob provides the account derived from his seeds. The compiler tries to sign the instruction with the computed seeds, resulting in a different PDA. The create account instruction is not signed by the account, and the instruction fails.

The PDA derived using the compiler's seed order is not a valid PDA, and the data account cannot be created using it. The contract needs to be redeployed with corrected annotations.

```
contract Test {  
  
    @payer(...)   
    @space(64)  
    @seed("A")  
    @bump("x")  
    @seed("B")  
    constructor() { [...] }  
    [...]  
}
```

*Figure 12.2: An example contract that is vulnerable to this issue*

### Recommendations

Short term, have the compiler raise an error if the developer places the bump annotation before a seed annotation. Alternatively, have the compiler place the bump value at the end of the seeds array irrespective of the annotation's position.

Long term, implement the compiler considering the expectations of the developer. Document the instances where the compiler diverges from these expectations.

### 13. Appending state variables to Solang contracts affects their storage layout

Severity: High

Difficulty: High

Type: Undefined Behavior

Finding ID: TOB-SOLCG-13

Target: <https://solang.readthedocs.io>

#### Description

Adding new state variables to a Solang-compiled contract will change the contract's storage layout. This behavior is different from Ethereum Solidity contracts. Developers who are not aware of the difference might brick their contracts by updating them with contracts containing additional state variables.

A Solang contract uses a linear byte array for storage. The byte array is divided into two sections, split by the offset `contract.fixed_layout_size`. The space from offset 0 to `fixed_layout_size` is used for storing fixed-size storage variables. The space from the `fixed_layout_size` index is considered to be a heap and is used for storing dynamic-size variables.

The `fixed_layout_size` offset depends on the contract's fixed-size state variables. The state variables are stored in the defined order. Appending new fixed-size variables would increase the `fixed_layout_size` offset, but the new variables will be stored from the previous `fixed_layout_size` offset.

As a result, if the contract is updated with a contract containing new fixed-size state variables, the new variables will be stored in the heap space of the old contract. This will corrupt the heap and result in an invalid state for the contract.

#### Exploit Scenario

Bob, the developer of a contract, adds new fixed-size state variables to the contract and updates the old contract using the new contract. Bob executes a function that writes to the first variable of the new state variables. The first variable is stored at the start of the heap of the old contract. The function overwrites the heap and corrupts the contract's state, and the contract becomes unusable.

#### Recommendations

Short term, add developer documentation to inform developers of these issues with updating to a contract with new state variables.

Long term, list the differences between Ethereum Solidity contracts and Solang contracts. Review the effects of these differences and document the issues stemming from them.

## 14. Fallback function does not verify the data account's magic value

Severity: High

Difficulty: Medium

Type: Data Validation

Finding ID: TOB-SOLCG-14

Target: `codegen/dispatch/solana.rs`

### Description

The storage for Solana contracts is represented using an account containing the contract's magic value. The fallback function does not check the magic value present in the account supplied as storage, thus allowing an attacker to send any account as storage and mislead the contract.

The code generator for the function dispatcher includes the magic value checks before calling the user-selected function (figure 14.1).

```
/// Add the dispatch for function given a matched selector
fn add_function_dispatch_case(
    [...]
) -> usize {

    let needs_account = if let ASTFunction::SolidityFunction(func_no) =
func_cfg.function_no {
        !ns.functions[func_no].is_pure()
    } else {
        true
    };

    if needs_account {
        // check for magic in data account, to see if data account is initialized
        let magic_ok = cfg.new_basic_block("magic_ok".into());
        let magic_bad = cfg.new_basic_block("magic_bad".into());
```

Figure 14.1: `codegen/dispatch/solana.rs#L279-L301`

The `add_function_dispatch_case` function handles the dispatch case for public functions. It adds the instructions to check the magic value and to revert the execution if it is incorrect.

The `function_dispatch` function handles the dispatch case for the fallback function (figure 14.2).

```
pub(crate) fn function_dispatch(
    [...]
```

```

) -> ControlFlowGraph {
  [...]
  match fallback {
    Some((cfg_no, _)) => {
      cfg.add(
        &mut vartab,
        Instr::Call {
          res: vec![],
          return_tys: vec![],
          args: vec![],
          call: InternalCallTy::Static { cfg_no },
        },
      );

      cfg.add(
        &mut vartab,
        Instr::ReturnCode {
          code: ReturnCode::Success,
        },
      );
    }
  }
}

```

Figure 14.2: [codegen/dispatch/solana.rs#L279-L301](#)

The `fallback` function is called without validating the magic value. If the `fallback` function is not pure and requires the data account, then the contract will be vulnerable. Because of the lack of the magic value check, an attacker can create their own account with data favorable to them and call the contract's `fallback` function. This might result in unintended behavior for the contract.

### Exploit Scenario

Alice, a developer, divides her contract into multiple sub-contracts and deploys them with different program IDs. Alice implements a proxy contract that contains a mapping from a function's selector to one of the sub-contracts' IDs. The proxy's `fallback` function uses the mapping to make a cross program invocation (CPI) call to the correct sub-contract for the given function ID.

One of the sub-contracts contains a helper function to call the token program and transfer the tokens. The tokens are owned by the PDA of the proxy contract. While calling the transfer helper function, the proxy signs the call with the PDA account.

Eve, an attacker, calls the proxy program with her data account. She sets the calldata to invoke the transfer helper function. The supplied data account contains Eve's exploit program for the transfer function's selector ID in the mapping. The proxy's `fallback` function calls Eve's program. As the proxy signs the call with the PDA, Eve's program transfers all tokens owned by the PDA to her account.

## Recommendations

Short term, add the `magic` value check to the `fallback` function's dispatcher if it is not pure.

Long term, list the validations that must be performed by a function and ensure that they are performed through comprehensive positive and negative tests.

## A. Vulnerability Categories

---

The following tables describe the vulnerability categories, severity levels, and difficulty levels used in this document.

Vulnerability Categories	
Category	Description
Access Controls	Insufficient authorization or assessment of rights
Auditing and Logging	Insufficient auditing of actions or logging of problems
Authentication	Improper identification of users
Configuration	Misconfigured servers, devices, or software components
Cryptography	A breach of system confidentiality or integrity
Data Exposure	Exposure of sensitive information
Data Validation	Improper reliance on the structure or values of data
Denial of Service	A system failure with an availability impact
Error Reporting	Insecure or insufficient reporting of error conditions
Patching	Use of an outdated software package or library
Session Management	Improper identification of authenticated users
Testing	Insufficient test methodology or test coverage
Timing	Race conditions or other order-of-operations flaws
Undefined Behavior	Undefined behavior triggered within the system

Severity Levels	
Severity	Description
Informational	The issue does not pose an immediate risk but is relevant to security best practices.
Undetermined	The extent of the risk was not determined during this engagement.
Low	The risk is small or is not one the client has indicated is important.
Medium	User information is at risk; exploitation could pose reputational, legal, or moderate financial risks.
High	The flaw could affect numerous users and have serious reputational, legal, or financial implications.

Difficulty Levels	
Difficulty	Description
Undetermined	The difficulty of exploitation was not determined during this engagement.
Low	The flaw is well known; public tools for its exploitation exist or can be scripted.
Medium	An attacker must write an exploit or will need in-depth knowledge of the system.
High	An attacker must have privileged access to the system, may need to know complex technical details, or must discover other weaknesses to exploit this issue.

## B. Non-Security-Related Findings

The following recommendations are not associated with specific vulnerabilities. However, implementing them will enhance code readability and may prevent the introduction of vulnerabilities in the future.

- **Reorganize the repository so that the root manifest is virtual (i.e., a workspace only).** Currently, the root manifest describes both a package and a workspace (figure B.1). The current organization complicates commands such as `cargo test`, as `--workspace` must be passed for the command to apply to the whole workspace, and not just the root package.

```
[package]
name = "solang"
...
[workspace]
members = ["solang-parser", "tests/wasm_host_attr"]
```

Figure B.1: *Cargo.toml#L1-L103*

- **Have the build script check that the correct version of `llvm-config` is referred to by `PATH`.** The expected version of `llvm-config` has the SBF target. The build script could run `llvm-config --targets-built` and verify that SBF appears in the output (figure B.2). Currently, if the wrong `llvm-config` is referred to by `PATH`, the build script will complete without error.

```
$ llvm-config --targets-built
AArch64 AMDGPU ARM AVR BPF Hexagon Lanai Mips MSP430 NVPTX PowerPC RISCV SBF
Sparc SystemZ VE WebAssembly X86 XCore
```

Figure B.2: *Output produced by the expected (patched) version of `llvm-config`*

- **Adopt a consistent import format.** (See figure B.3.) Doing so will make it easier to determine what symbols are imported and from where. `Rustfmt`'s (unstable) `imports_granularity` and `group_imports` configurations could help with this.

```
use self::{
    cfg::{optimize_and_check_cfg, ControlFlowGraph, Instr},
    dispatch::function_dispatch,
    expression::expression,
    solana_accounts::account_collection::collect_accounts_from_contract,
    variable::Vartable,
};
...
use crate::codegen::cfg::ASTFunction;
```

Figure B.3: *codegen/mod.rs#L27-L43*



- **Run Clippy's pedantic lints in the CI process.** As previously reported (in finding TOB-SOLANG-3 in the "Solang Parser and Semantic Analysis" report), Clippy's pedantic lints produce many warnings when applied to the codebase. Addressing them would improve the quality of the code. Example warnings appear in figures B.4 through B.7.

```
warning: redundant closure
--> src/codegen/cfg.rs:1802:14
|
1802 |         .map(|stmt| stmt.reachable())
|               ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^ help: replace the closure with the
method itself: `sema::ast::Statement::reachable`
|
= help: for further information visit
https://rust-lang.github.io/rust-clippy/master/index.html#redundant\_closure\_for\_method\_calls
```

*Figure B.4: The warning produced by redundant\_closure\_for\_method\_calls*

```
warning: implicitly cloning a `Vec` by calling `to_vec` on its dereferenced
type
--> src/codegen/expression.rs:2519:16
|
2519 |         value: id.to_vec(),
|               ^^^^^^^^^^^^^ help: consider using: `id.clone()`
|
= help: for further information visit
https://rust-lang.github.io/rust-clippy/master/index.html#implicit\_clone
```

*Figure B.5: The warning produced by implicit\_clone*

```
warning: used `cloned` where `copied` could be used instead
--> src/codegen/solana_accounts/account_management.rs:21:14
|
21 |         .cloned()
|         ^^^^^^^ help: try: `copied`
|
= help: for further information visit
https://rust-lang.github.io/rust-clippy/master/index.html#cloned\_instead\_of\_copied
```

*Figure B.6: The warning produced by cloned\_instead\_of\_copied*

```
warning: it is more concise to loop over containers instead of using explicit
iteration methods
--> src/codegen/strength_reduce/mod.rs:91:29
|
91 |         for (block_no, vars) in block_vars.into_iter() {
|                                   ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^ help: to write this
more concisely, try: `block_vars`
|
```

```

= help: for further information visit
https://rust-lang.github.io/rust-clippy/master/index.html#explicit\_into\_iter\_loop

```

Figure B.7: The warning produced by `explicit_into_iter_loop`

- **Eliminate the unnecessary use of `mut` in figure B.8.**

```

warning: variable does not need to be mutable
--> src/codegen/subexpression\_elimination/mod.rs:165:13
|
165 |         let mut cur_block = &mut cfg.blocks[*block_no];
|         -----^^^^^^^^^^^^^^
|         |
|         help: remove this `mut`
|
= note: `#[warn(unused_mut)]` on by default

```

Figure B.8: The warning produced by `unused_mut`

- **Change the use of `borrow_mut` to `borrow` in figure B.9.** Using `borrow_mut` unnecessarily could result in a panic. (Note that `unnecessary_borrow_mut` is a Dylint lint.)

```

warning: borrowed reference is used only immutably
-->
src/codegen/subexpression\_elimination/available\_expression\_set.rs:368:43
|
368 |         for (child_id, node) in &var_node.borrow_mut().children {
|                                     ^^^^^^^^^^^^^^^^^ help: use:
`borrow()`
|
= note: `#[warn(unnecessary_borrow_mut)]` on by default

```

Figure B.9: The warning produced by `unnecessary_borrow_mut`

- **Eliminate the unnecessary call to `as_bytes` in figure B.10.** (Note that `unnecessary_conversion_for_trait` is a Dylint lint.)

```

warning: the receiver implements the required traits
--> src/codegen/events/solana.rs:34:23
|
34 |         hasher.update(discriminator_image.as_bytes());
|                                     ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^ help: use:
`&discriminator_image`

```

Figure B.10: The warning produced by `unnecessary_conversion_for_trait`

- **Eliminate the duplicate dependencies that appear in the root manifest.** The packages sha2 and tempfile appear as both regular and “dev” dependencies. It is sufficient that they appear as just regular dependencies.

```
[dependencies]
...
tempfile = "3.4"
...
sha2 = "0.10"
...
[dev-dependencies]
...
sha2 = "0.10"
...
tempfile = "3.3"
```

Figure B.11: *Cargo.toml#L18–L85*

- **Eliminate the corner case that can cause test\_mul\_within\_range\_signed to fail.** If first\_operand\_rand is  $-2^{(N-1)}$  and second\_op is  $-1$ , the multiplication operation will overflow.

```
#[test]
fn test_mul_within_range_signed() {
    let mut rng = rand::thread_rng();
    for width in (8..=256).step_by(8) {
        ...
        // The range of values that can be held in signed N bits is  $[-2^{(N-1)}, 2^{(N-1)}-1]$ . Here we generate a random number within this range and multiply it by -1, 1 or 0.
        let first_operand_rand = rng.gen_bigint(width - 1).sub(1_u32);
        println!("First op : {first_operand_rand:?}");

        let side = vec![-1, 0, 1];
        // -1, 1 or 0
        let second_op = BigInt::from(*side.choose(&mut rng).unwrap());
        println!("second op : {second_op:?}");
    }
}
```

Figure B.12: *tests/solana\_tests/primitives.rs#L989–L1011*

- **Replace the call to `BigUint::pow` followed by `truncate_biguint` (figure B.13) with just one call to `BigUint::modpow` (figure B.14).** Doing so will make the uint test more efficient.

```
let mut res = a.clone().pow(n);
truncate_biguint(&mut res, width);
```

Figure B.13: *tests/solana\_tests/primitives.rs#L543–L544*

```
let res = a
```

```

        .clone()
        .modpow(&BigUint::from(n), &BigUint::from(2u64).pow(width as u32));
    // truncate_biguint(&mut res, width);

```

Figure B.14: The proposed change to the code in figure B.13

- **Check both sides of the boundary condition in the code in figure B.15 (i.e., add code like in figure B.16 to transfer\_fails\_not\_enough).** Doing so will help increase confidence in the transfer\_fails\_not\_enough test.

```

let res = vm.function_must_fail(
    "transfer",
    &[
        BorshToken::FixedBytes(new.to_vec()),
        BorshToken::Uint {
            width: 64,
            value: BigInt::from(104u8),
        },
    ],
);
assert!(res.is_err());

```

Figure B.15: [tests/solana\\_tests/balance.rs#L256-L266](#)

```

let res = vm.function_must_fail(
    "transfer",
    &[
        BorshToken::FixedBytes(new.to_vec()),
        BorshToken::Uint {
            width: 64,
            value: BigInt::from(103u8),
        },
    ],
);
assert!(res.is_ok());

```

Figure B.16: The proposed change to the code in figure B.13

A similar recommendation applies to the transfer\_fails\_overflow test. (See figure B.17).

```

let res = vm.function_must_fail(
    "transfer",
    &[
        BorshToken::FixedBytes(new.to_vec()),
        BorshToken::Uint {
            width: 64,
            value: BigInt::from(104u8),
        },
    ],
);

```

```
assert!(res.is_err());
```

Figure B.17: [tests/solana\\_tests/balance.rs#L297-L307](#)

- **Rename the following methods to better communicate what they do:**
  - Rename `function_must_fail` to `function_may_fail`.
  - Rename `edges` to `successors`.
  - Rename `clone_for_parent_block` to `deep_clone`.

- **Correct the grammar in the comments in figures B.18 and B.19.**

```
/// When a reaching definition change, we remove the variable node and all its  
descendants from  
/// the graph
```

Figure B.18: The word `change` should likely be `changes`.

([codegen/subexpression\\_elimination/available\\_expression\\_set.rs#L358-L359](#))

```
/// Regenerate instructions after that we exchanged common subexpressions for  
temporaries
```

Figure B.19: The word `exchanged` should likely be `exchange`.

([codegen/subexpression\\_elimination/instruction.rs#L203](#))

- **Swap the comments in figure B.20, which appear to be associated with the wrong functions.**

```
/// Get the maximum unsigned value in a set  
pub(super) fn set_max_signed(set: &HashSet<Value>) -> Option<BigInt> {  
    ...  
    /// Get the maximum signed value in a set  
    pub(super) fn set_max_unsigned(set: &HashSet<Value>) -> BigInt {
```

Figure B.20: [codegen/strength\\_reduce/value.rs#L69-L95](#)

- **Correct the typo in `expresson_known_bits` (figure B.21).**

```
fn expresson_known_bits() {
```

Figure B.21: The word `expresson` should be `expression`.

([codegen/strength\\_reduce/tests.rs#L30](#))

- **Rewrite the code in figure B.22 to use `unwrap` or `expect`.** Doing so will make the code more clear.

```

if let Some(block_vars) = block_vars.get_mut(&edge) {
    ...
} else {
    unreachable!();
}

```

Figure B.22: *codegen/dead\_storage.rs#L149–L171*

- **Use named constants in place of magic numbers throughout the code.** Doing so will make the code more clear. Examples where magic numbers are used appear in figures B.23 through B.25.

```

let lamports_runtime_constant = (128 + space_runtime_constant) * 3480 * 2;

```

Figure B.23: *codegen/solana\_deploy.rs#L342*

```

flow[block_1] = BigRational::from_integer(1000.into());

```

Figure B.24:

*codegen/subexpression\_elimination/anticipated\_expressions.rs#L118*

```

&& BigRational::from_integer(2000.into()) == *flow_magnitude

```

Figure B.25:

*codegen/subexpression\_elimination/anticipated\_expressions.rs#L161*

In some cases, even replacing 0 with a named constant would make the code more clear. For example, in figure B.26, 0 could be replaced with ENTRY\_BLOCK.

```

vars[0].clone()

```

Figure B.26: *codegen/dead\_storage.rs#L114*

- **Add a comment explaining why it is acceptable that the `highest_set_bit` function (figure B.27) returns 0 for both 0 and 1.** While this behavior does not appear to cause a problem now, the function could easily be misused in future code.

```

fn highest_set_bit(bs: &[u8]) -> usize {
    for (i, b) in bs.iter().enumerate().rev() {
        if *b != 0 {
            return (i + 1) * 8 - bs[i].leading_zeros() as usize - 1;
        }
    }
    0
}

```

Figure B.27: *codegen/strength\_reduce/mod.rs#L569–L577*

- **Remove code to handle the receive function in the Solana function dispatch generator (figure B.28).** The Solang compiler does not support the receive function for the Solana target. The compiler will raise an error if a contract contains a receive function. As a result, it is unnecessary to check for the presence of the receive function in the dispatch generator. Removing the unnecessary code will reduce the code complexity and increase readability.

```
let receive = all_cfg
    .iter()
    .enumerate()
    .find(|(_, cfg)| cfg.public && cfg.ty == pt::FunctionTy::Receive);

if fallback.is_none() && receive.is_none() {
```

Figure B.28: *codegen/dispatch/solana.rs#L227-L232*